



# 데이터 과학 외전

Day 4 – 앙상블 알고리즘

# Contents

---


- Bias and Variance
- Bagging and Boosting
- Random Forest
- Adaboost and XGBoost

# Bias And Variance



사람들의 나이와 연간의료비용을 데이터로 나타내었다

- 나이가 어리거나 젊을 때는 의료비용의 증가가 없거나 거의 미미하다
- 중년 이상이 되면서 의료비용이 서서히 증가하기 시작한다
- 노년 이상이 되면서 의료비용의 증가 폭이 더욱 커진다

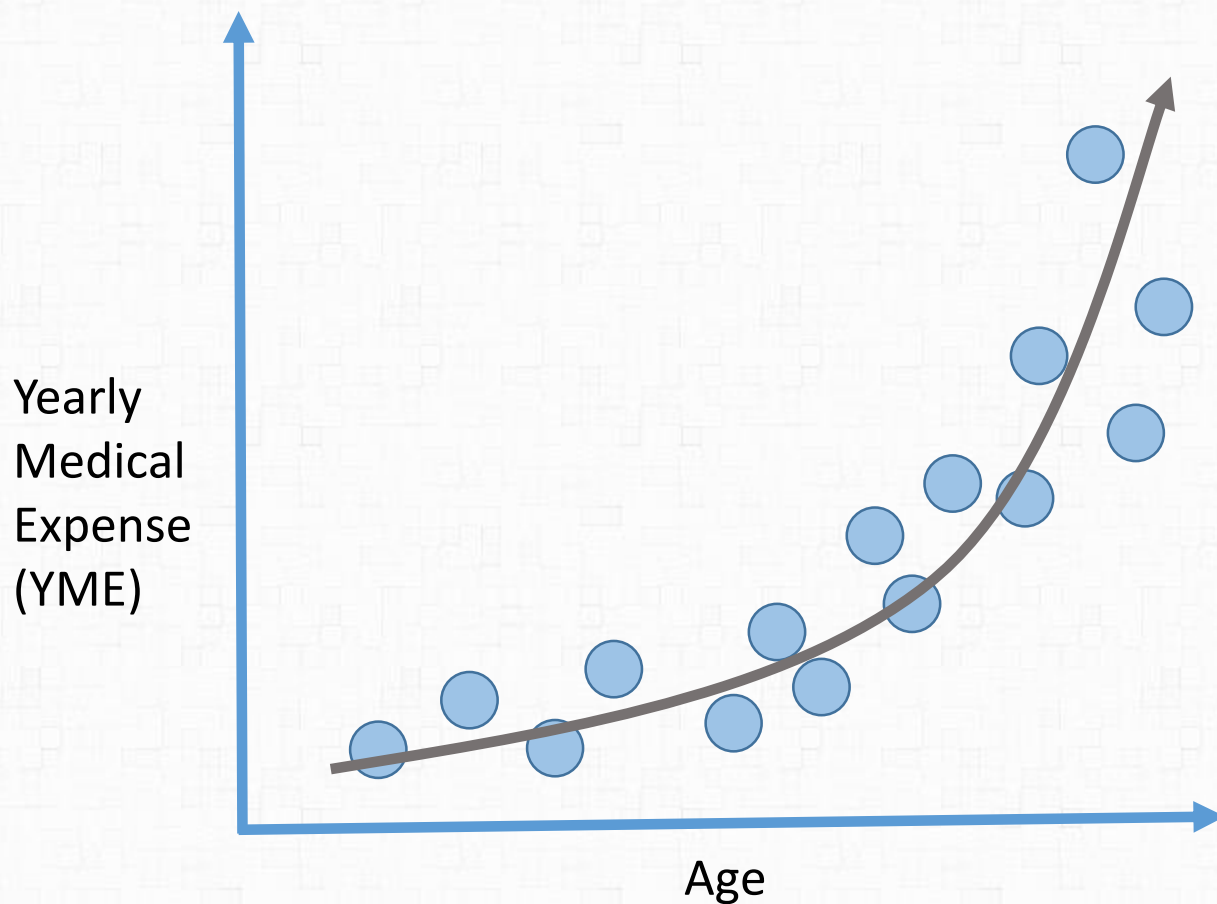


이상적으로는

의료비용과 나이의 관계를 결정하는  
법칙 = 수식 = 모델  
을 알고자 함


“True” relationship

나이와 의료비용의 관계를 sample data로  
부터 추정하는 과정 = data modeling



데이터를 모델을 학습(추정)하기 위한 데이터(**학습데이터**)

모델을 검증하기 위한 데이터(**테스트 데이터**)로 구분한다




## First Try: Linear Regression

$$y = \alpha_0 + \alpha_1 x$$

y : YME

x : Age




## First Try: Linear Regression

$$y = \alpha_0 + \alpha_1 x$$

y : YME

x : Age






## First Try: Linear Regression

$$y = \alpha_0 + \alpha_1 x$$

y : YME

x : Age



## First Try: Linear Regression



$$y = \alpha_0 + \alpha_1 x$$

y : YME

x : Age

**Linear Regression  
will never capture  
the “true relationship”**

**= Bias**




## Second Try: Polynomial Regression

**Model is flexible to capture the non-linearity**

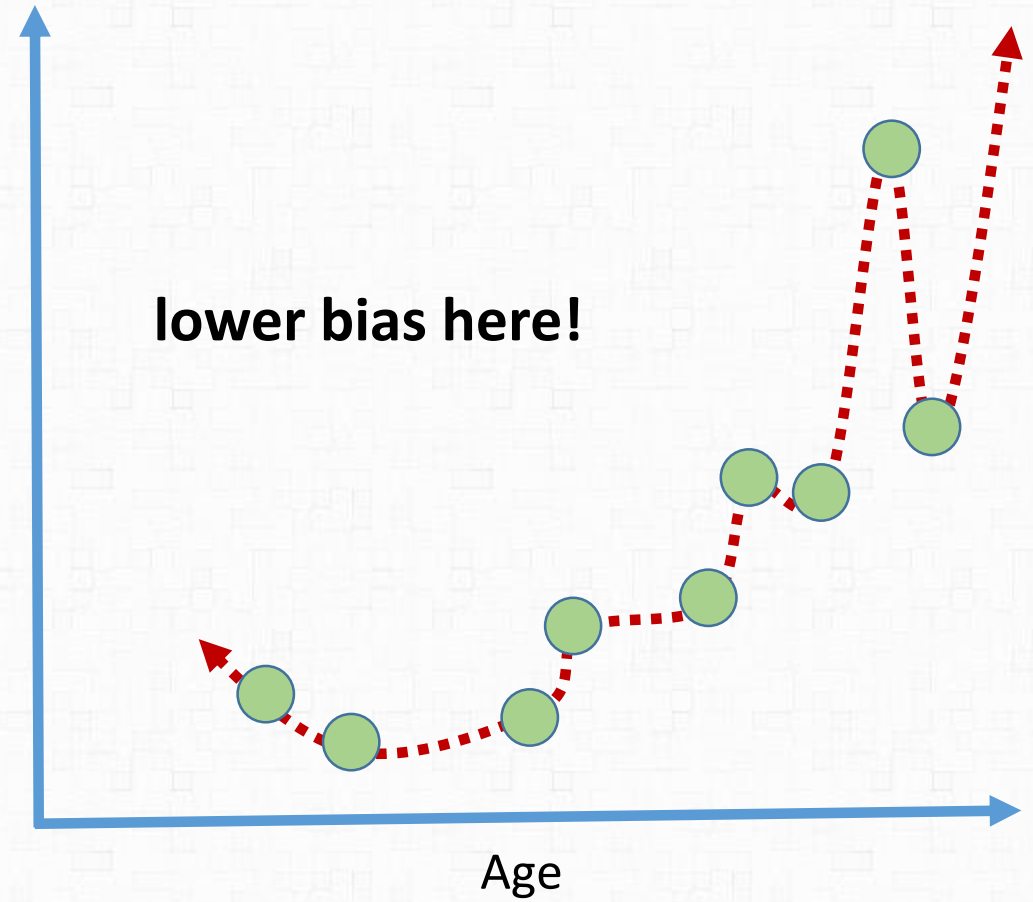

$$y = \alpha_0 + \alpha_1 x + \alpha_2 x^2 + \alpha_3 x^3 + \alpha_4 x^4 + \alpha_5 x^5$$

y : YME

x : Age





## Second Try: Polynomial Regression





## Second Try: Polynomial Regression

**what about training set**

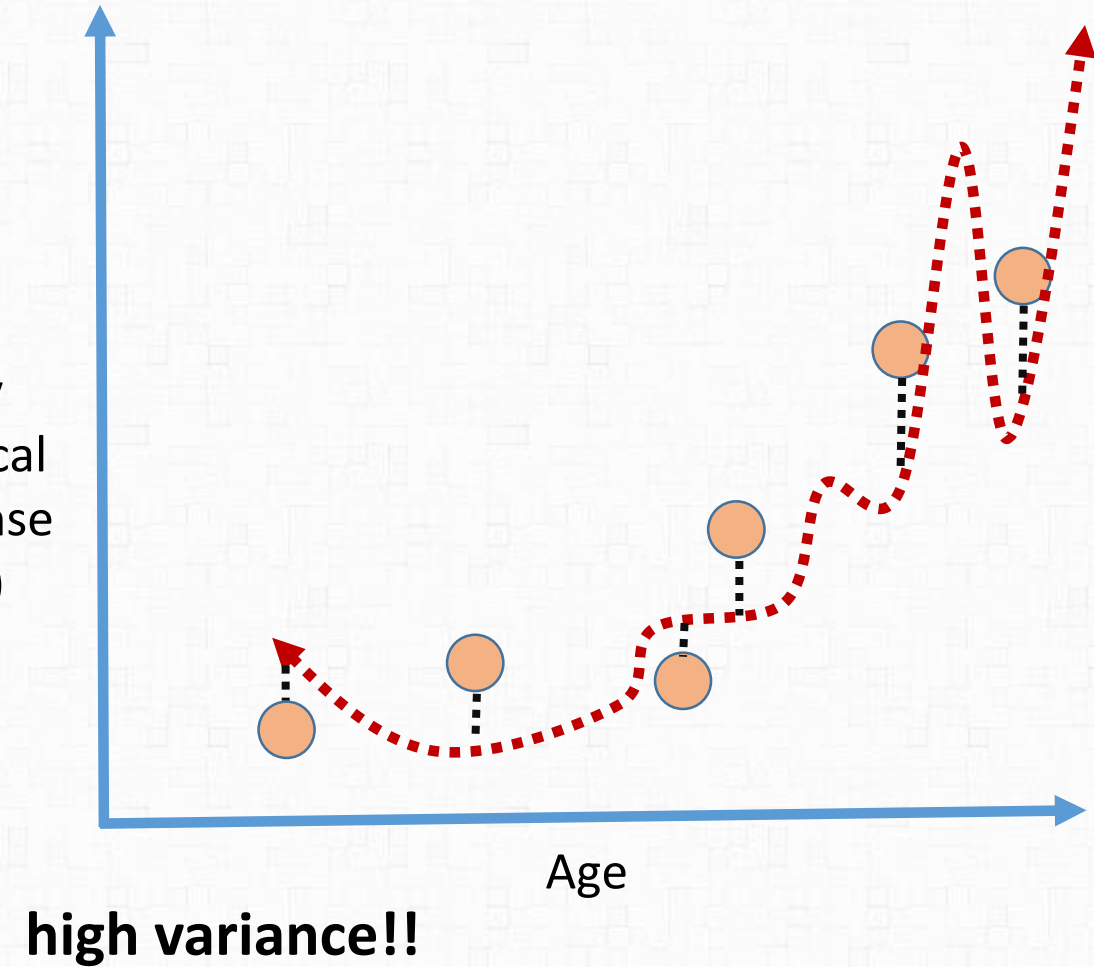



## Second Try: Polynomial Regression

**what about training set**




The difference in fit between two datasets is call **Variance**






low variance model consistently give a good prediction  
(even though it is not a great prediction)

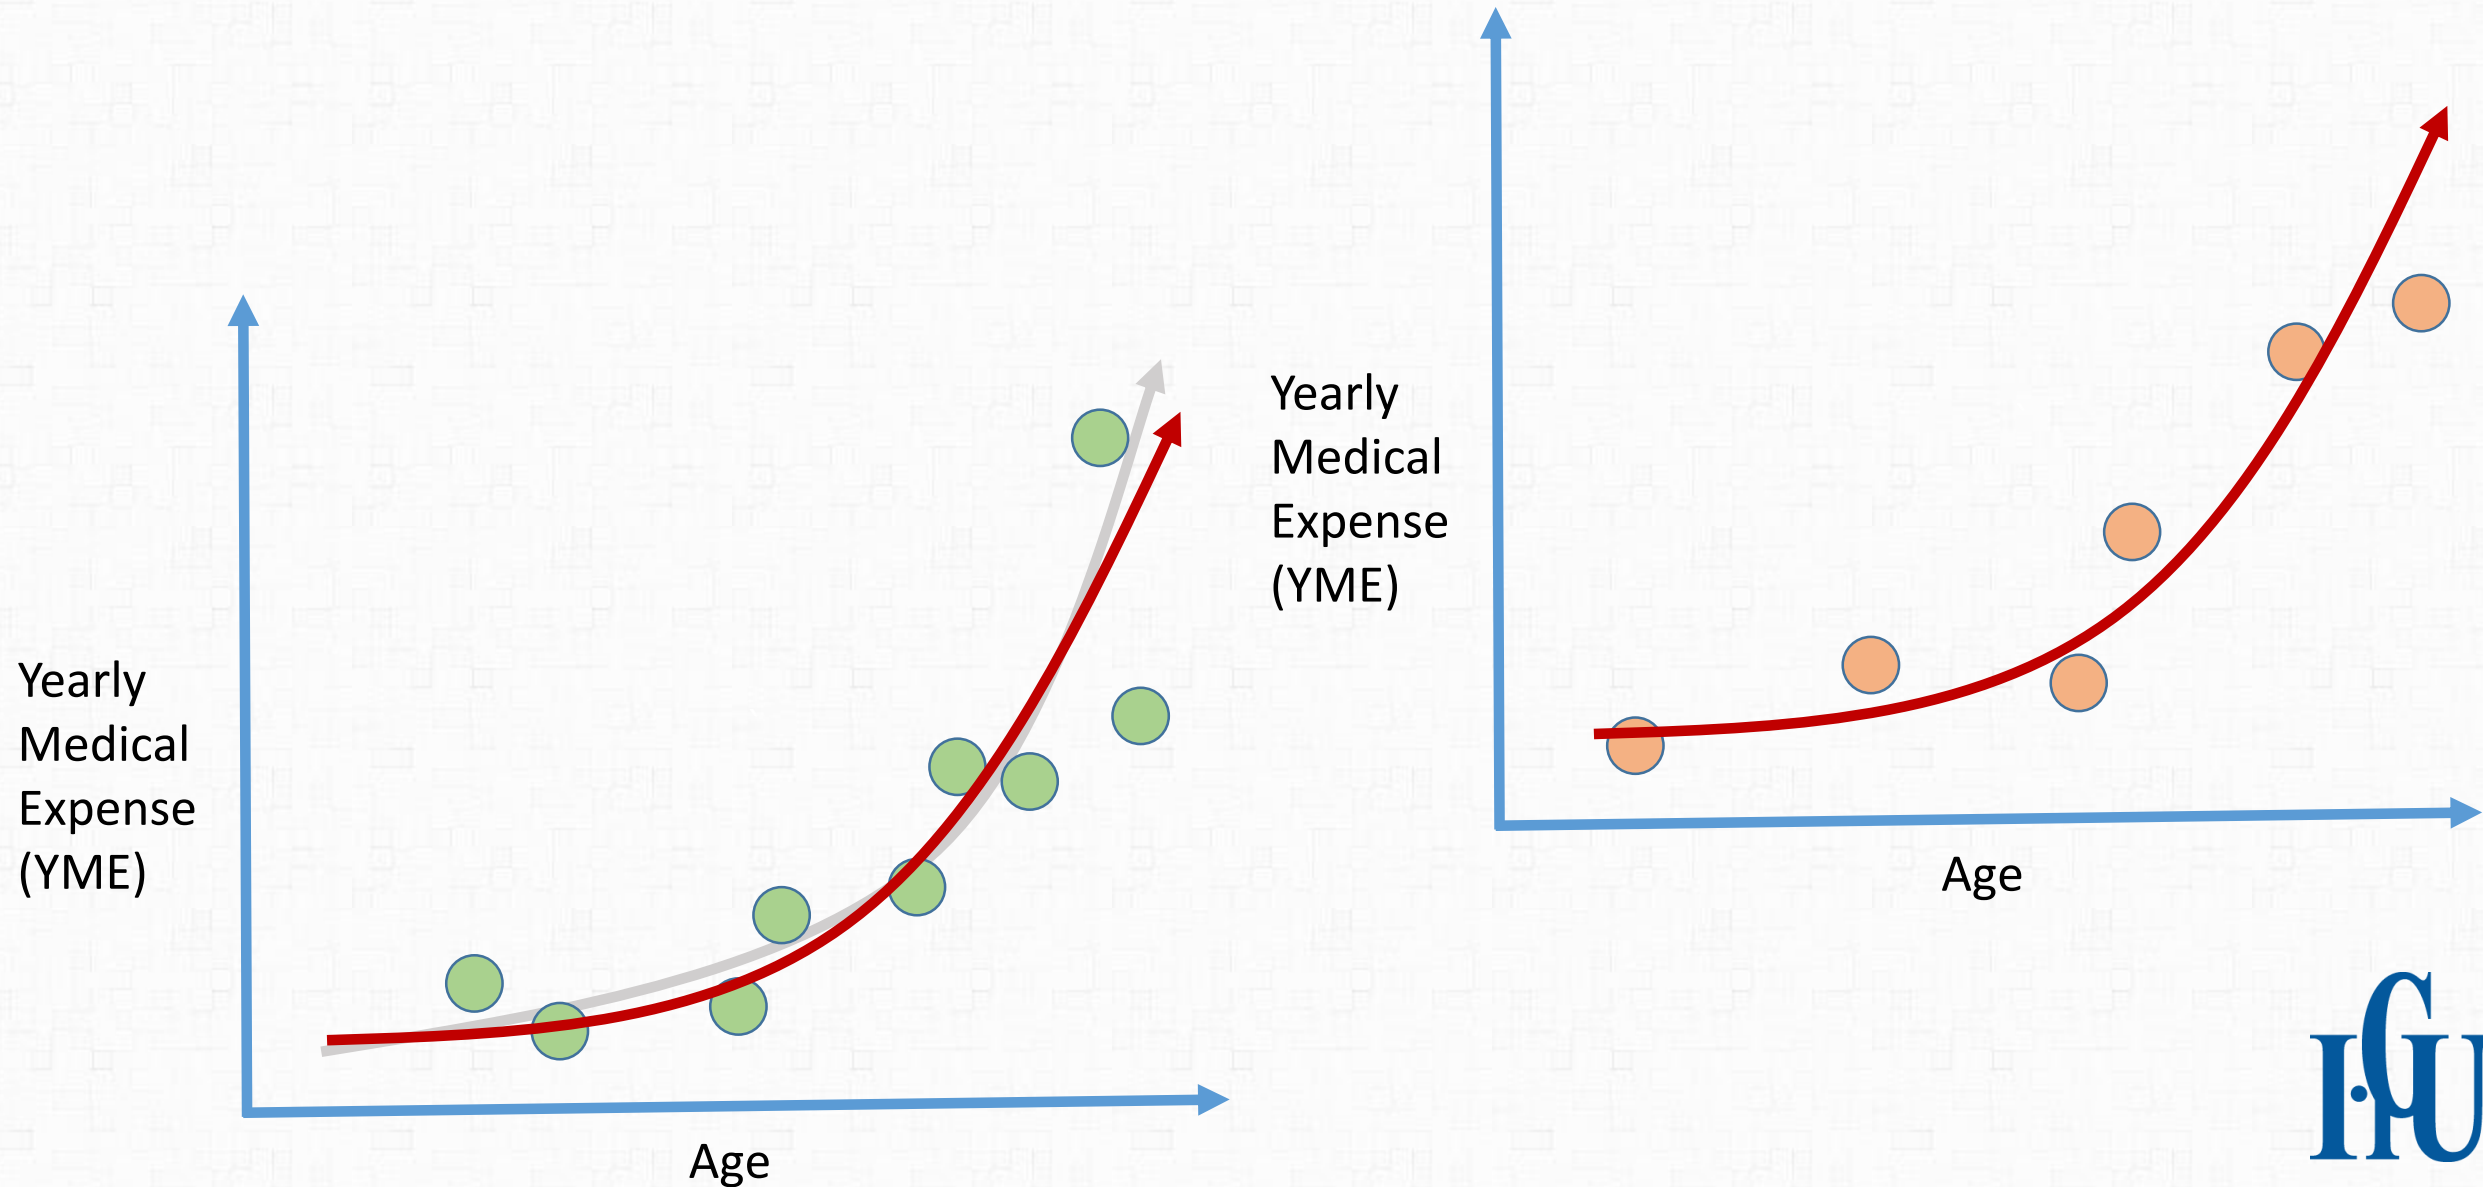




the ideal model has **low bias**,  
which accurately model the true relationship





... and it has **low variance**,  
by making consistent predictions across different datasets



The key is to find a balance between a simple model and a complex model

**The commonly used methods are:  
regularization, boosting and bagging**



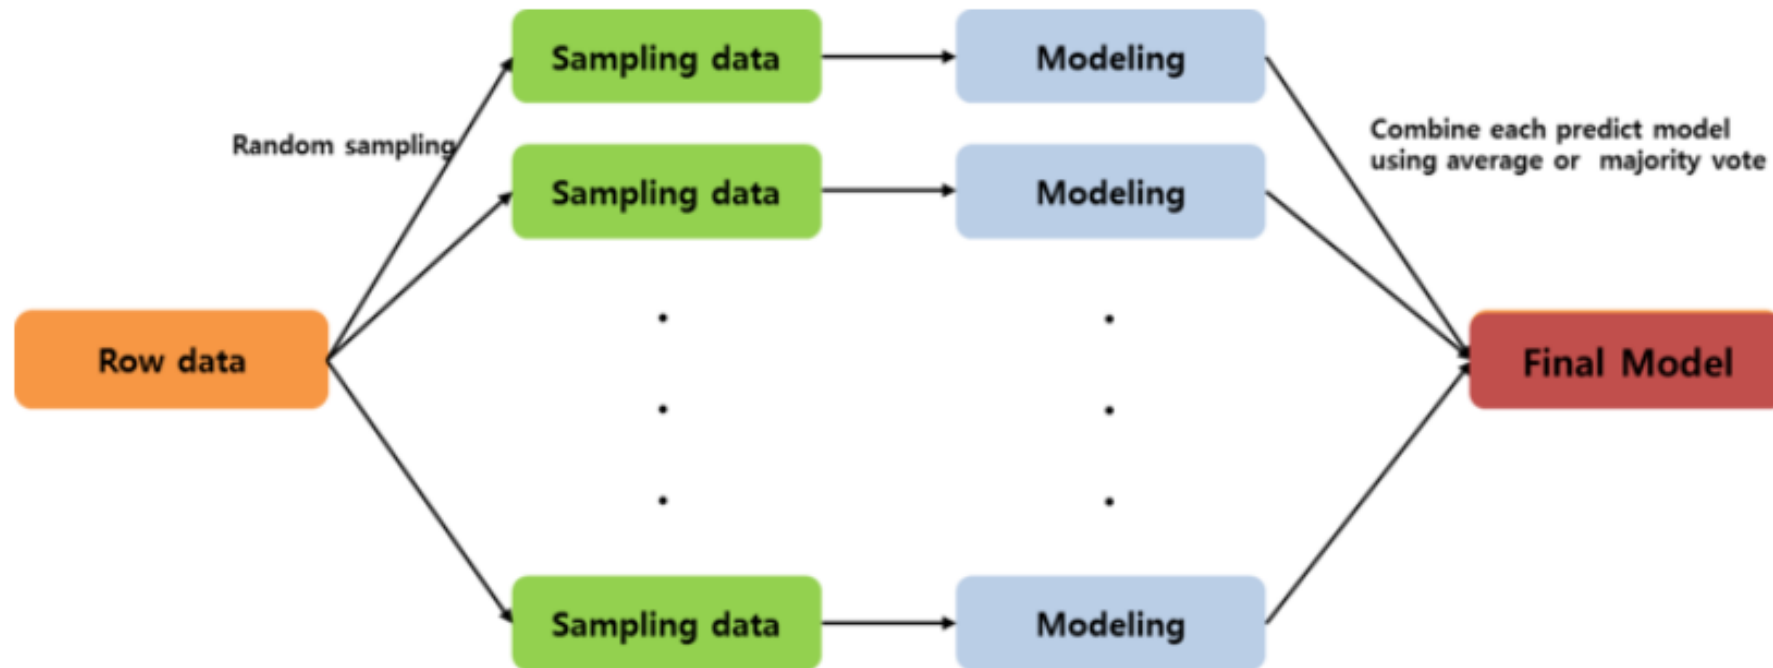
# Bagging and Boosting

---

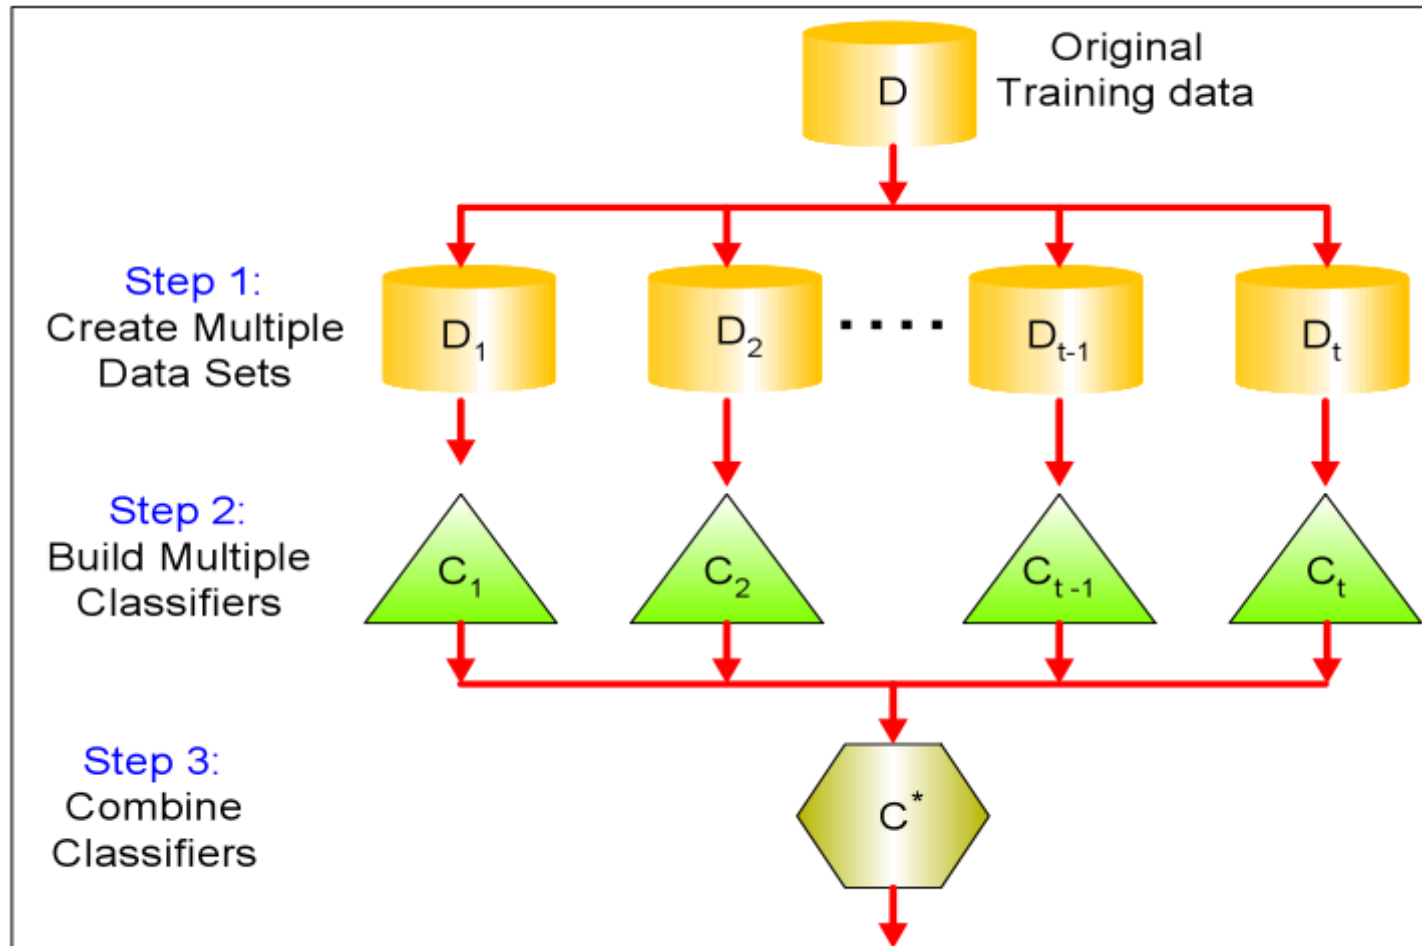
---

---

# Bagging



# Bagging



# Example of Bagging

- Sampling with replacement


Training Data

Data ID	1	2	3	4	5	6	7	8	9	10
Original Data	1	2	3	4	5	6	7	8	9	10
Bagging (Round 1)	7	8	10	8	2	5	10	10	5	9
Bagging (Round 2)	1	4	9	1	2	3	2	7	3	2
Bagging (Round 3)	1	8	5	10	5	5	9	6	3	7

- Build classifier on each bootstrap sample
- Each data point has probability  $(1 - 1/n)^n$  of being selected as test data
- Training data =  $1 - (1 - 1/n)^n$  of the original data

# Example of Bagging


---






# Example of Bagging

---



# 100 bagged trees



shades of blue/red indicate strength of vote for particular classification

# Bagging / Resampling

---

- It reduces variance of model complexity
- It builds various models with many different samplings
- It reduces bias b/c it combines all models
- Aggregating all the models, reduce variance

# Practice - Bagging

```
spamD <- read.table('spamD.tsv',header=T,sep='\t')
```

```
spamD[1:5, c(1:5,58:59)]
```

```
##      word.freq.make word.freq.address word.freq.all word.freq.3d
## 1           0.00           0.64           0.64           0
## 2           0.21           0.28           0.50           0
## 3           0.06           0.00           0.71           0
## 4           0.00           0.00           0.00           0
## 5           0.00           0.00           0.00           0
```

```
##      word.freq.our spam rgroup
## 1           0.32 spam      52
## 2           0.14 spam      91
## 3           1.23 spam      49
## 4           0.63 spam      88
## 5           0.63 spam      73
```

```
spamTrain <- subset(spamD,spamD$rgroup>=10)
spamTest  <- subset(spamD,spamD$rgroup<10)
```

```
# Note 1:
#   Load the data and split into training (90% of data)
#   and test (10% of data) sets.
```

# Building Tree model

```
# Note 2:
#   Use all the features and do binary classification,
#   where TRUE corresponds to spam documents.
```

```
spamVars <- setdiff(colnames(spamD), list('rgroup', 'spam'))
spamFormula <- as.formula(paste('spam=="spam"',
                                paste(spamVars, collapse=' + '), sep=' ~ '))
```

```
# Note 4:
#   A function to calculate and return various measures
#   on the model: prediction accuracy, and f1, which is the
#   harmonic mean of precision and recall.
```

```
accuracyMeasures <- function(pred, truth, name="model") {
  ctable <- table(truth=truth,
                  pred=(pred>0.5))
  accuracy <- sum(diag(ctable))/sum(ctable)
  precision <- ctable[2,2]/sum(ctable[,2])
  recall <- ctable[2,2]/sum(ctable[2,])
  f1 <- 2*precision*recall/(precision+recall)
  data.frame(model=name, accuracy=accuracy, f1=f1)
}
```

```
# Note 6:
#   Convert the class probability estimator into a
#   classifier by labeling documents that score greater
#   than 0.5 as
#   spam.
```

```
library(rpart)
treemodel <- rpart(spamFormula, spamTrain)

accuracyMeasures(predict(treemodel, newdata=spamTrain),
                  spamTrain$spam=="spam",
                  name="tree, training")

##           model  accuracy      f1
## 1 tree, training 0.9104514 0.88337

accuracyMeasures(predict(treemodel, newdata=spamTest),
                  spamTest$spam=="spam",
                  name="tree, test")

##           model  accuracy      f1
## 1 tree, test 0.8799127 0.8414986
```

## bagging

```
ntrain <- dim(spamTrain)[1]  
n <- ntrain  
ntree <- 100
```

```
# Note 1:  
#   Use bootstrap samples the same size as the training  
#   set, with 100 trees.
```

```
samples <- sapply(1:ntree,  
                 FUN = function(iter)  
                 {sample(1:ntrain, size=n, replace=T)})
```

```
# Note 2:  
#   Build the bootstrap samples by sampling the row indices  
#   of spamTrain with replacement. Each  
#   column of the matrix samples represents the row indices  
#   into spamTrain  
#   that comprise the bootstrap sample.
```

```
treelist <- lapply(1:ntree,  
                 FUN=function(iter)  
                 {samp <- samples[,iter];  
                  rpart(spamFormula, spamTrain[samp,])})
```

```
# Note 3:  
#   Train the individual decision trees and return them  
#   in a list. Note: this step can take a few minutes.
```

```

predict.bag <- function(treelist, newdata) {
  preds <- sapply(1:length(treelist),
                 FUN=function(iter) {
                   predict(treelist[[iter]], newdata=newdata)})
  predsums <- rowSums(preds)
  predsums/length(treelist)
}

```

# Note 4:  
 # predict.bag assumes the underlying classifier returns  
 # decision probabilities, not  
 # decisions.

```

accuracyMeasures(predict.bag(treelist, newdata=spamTrain),
                  spamTrain$spam=="spam",
                  name="bagging, training")

```

```

##           model  accuracy          f1
## 1 bagging, training 0.9227613 0.8990536

```

```

accuracyMeasures(predict.bag(treelist, newdata=spamTest),
                  spamTest$spam=="spam",
                  name="bagging, test")

```

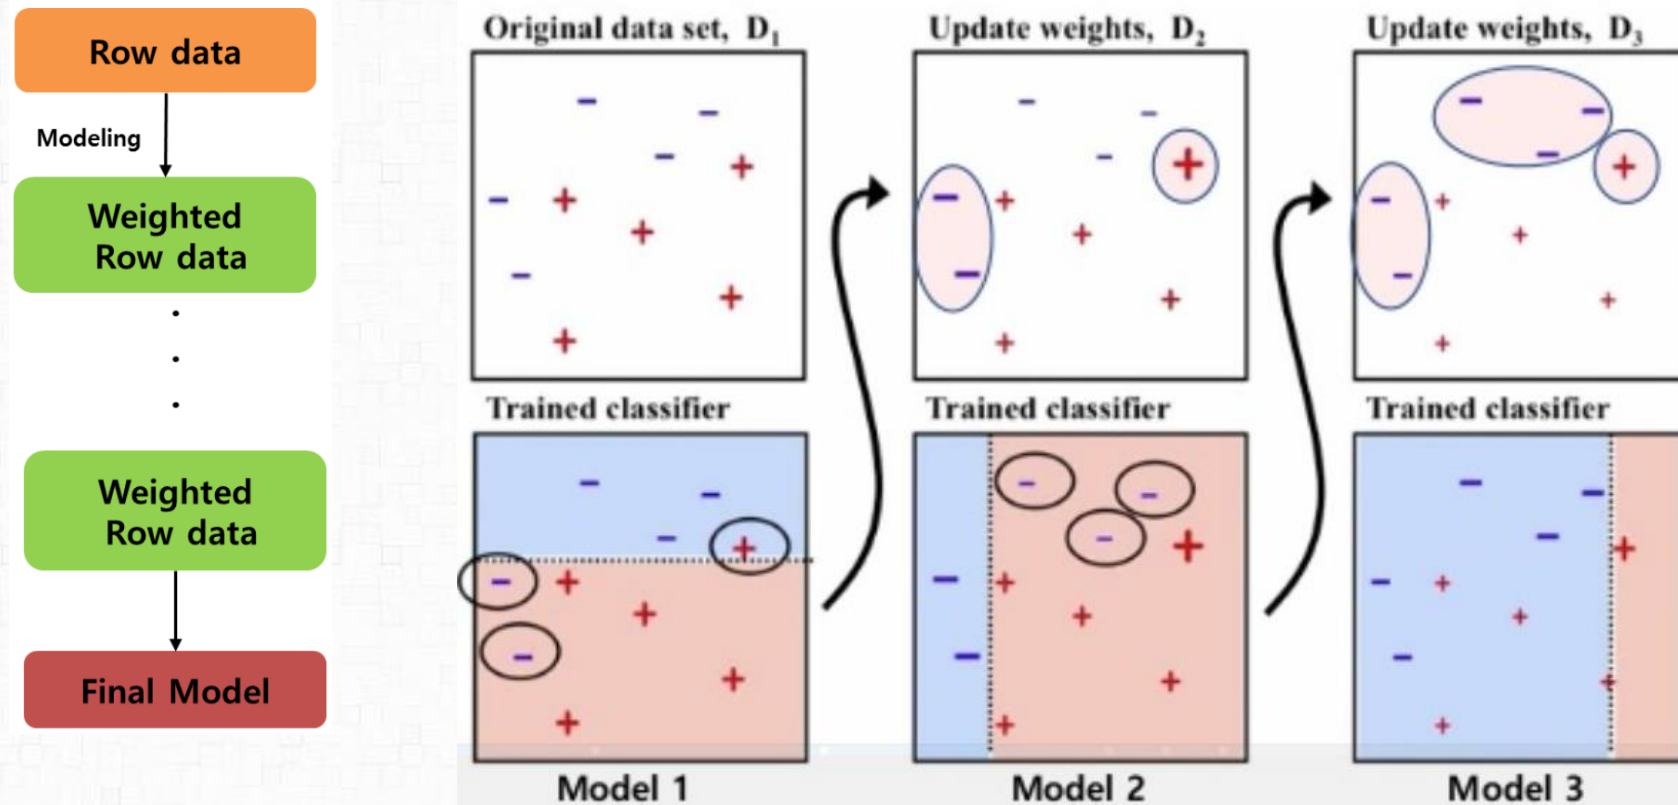
```

##           model  accuracy          f1
## 1 bagging, test 0.9126638 0.8809524

```



# Boosting (C5.0)



# Boosting

- Aggregate the result of boosted model

알고리즘	특징	비고
AdaBoost	<ul style="list-style-type: none"> <li>• 다수결을 통한 정답 분류 및 오답에 가중치 부여</li> </ul>	
GBM	<ul style="list-style-type: none"> <li>• Loss Function의 gradient를 통해 오답에 가중치 부여</li> </ul>	<a href="#">gradient_boosting.pdf</a>
Xgboost	<ul style="list-style-type: none"> <li>• GBM 대비 성능향상</li> <li>• 시스템 자원 효율적 활용 ( CPU, Mem)</li> <li>• Kaggle을 통한 성능 검증 (많은 상위 랭커가 사용)</li> </ul>	2014년 공개 <a href="#">boosting-algorithm-xgboost</a>
Light GBM	<ul style="list-style-type: none"> <li>• Xgboost 대비 성능향상 및 자원소모 최소화</li> <li>• Xgboost가 처리하지 못하는 대용량 데이터 학습 가능</li> <li>• Approximates the split (근사치의 분할)을 통한 성능 향상</li> </ul>	2016년 공개 <a href="#">light-gbm-vs-xgboost</a>

# Bagging v.s. Boosting

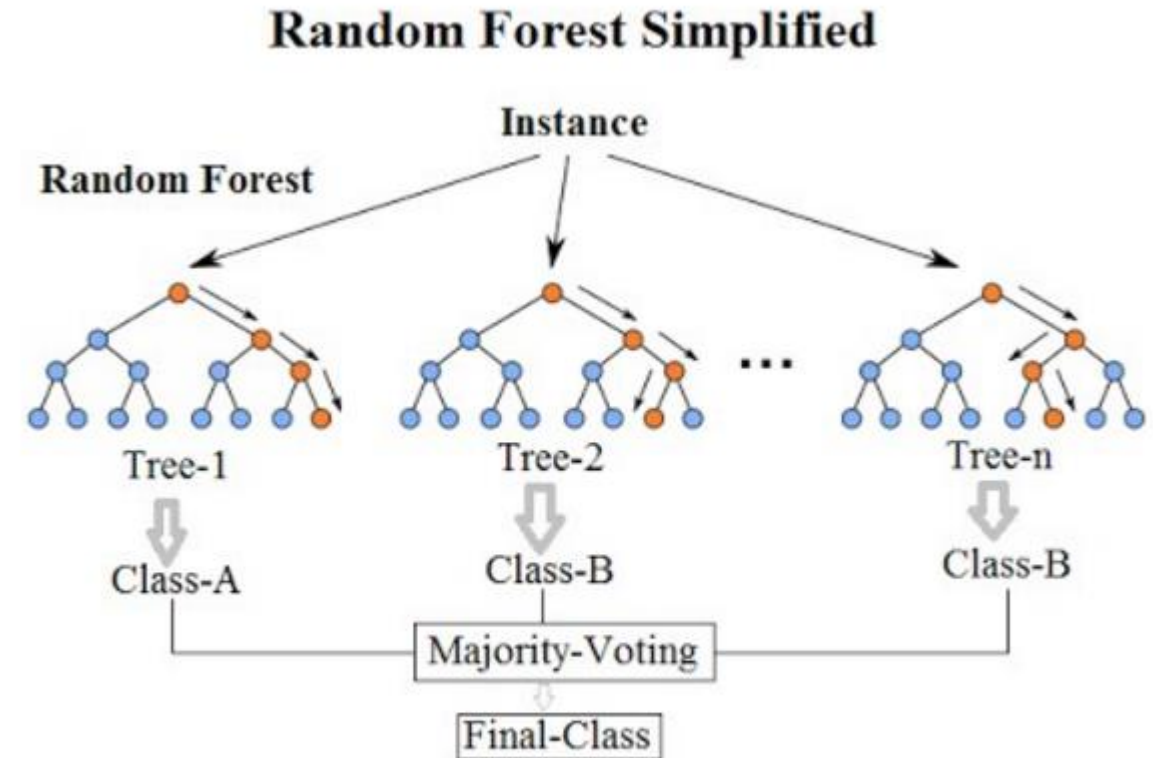
비교	Bagging	Boosting
특징	병렬 앙상블 모델 (각 모델은 서로 독립적)	연속 앙상블 (이전 모델의 오류를 고려)
목적	Variance 감소	Bias 감소
적합한 상황	복잡한 모델 (High variance, Low bias)	Low variance, High bias 모델
대표 알고리즘	Random Forest	Gradient Boosting, AdaBoost
Sampling	Random Sampling	Random Sampling with weight on error

# Random Forest




# Random Forest

- Tree-based Ensemble model with Bagging



# Random Forest – Tree Correlation

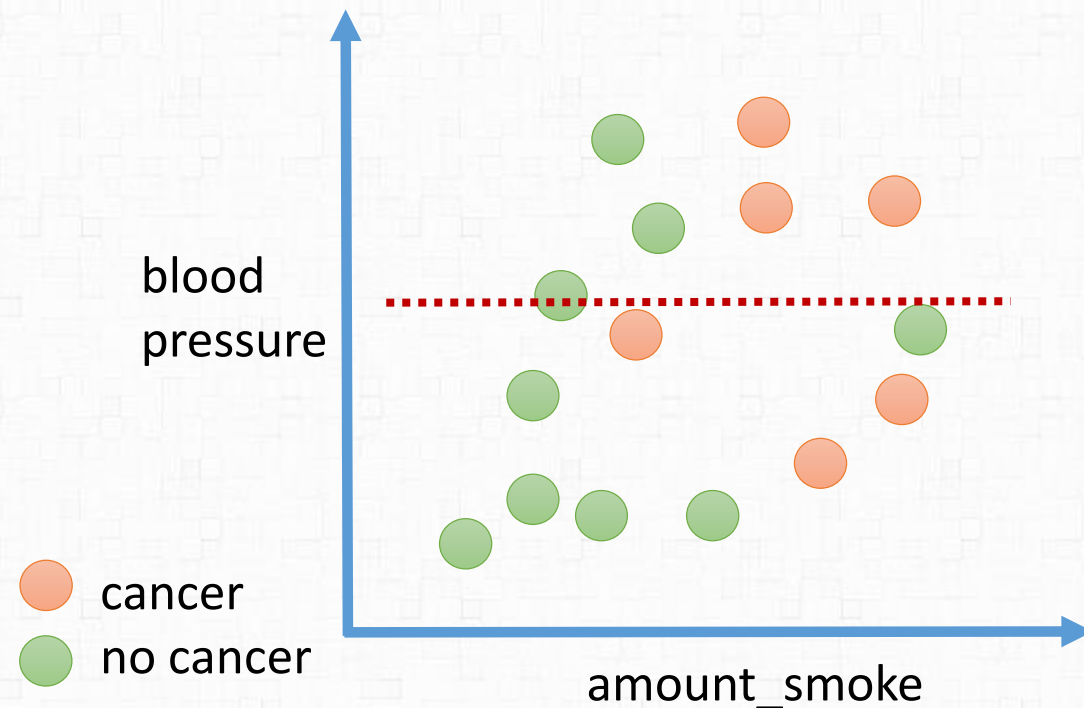
- Bagging reduces variance by aggregating many different models
- Bagging samples **observations** not features
- If certain features have high predictive power then all the tree models look very similar
  - Eg. Smoker variable of insurance dataset
  - High correlation in trees
- Random Forest
  - Select **features** randomly and build trees





# Random Forest – Tree Correlation

- Bagging reduces variance by aggregating many different models
- Bagging samples **observations** not features
- If certain features have high predictive power then all the tree models look very similar
  - Eg. Smoker variable of insurance dataset
  - High correlation in trees
- Random Forest
  - Select **features** randomly and build trees



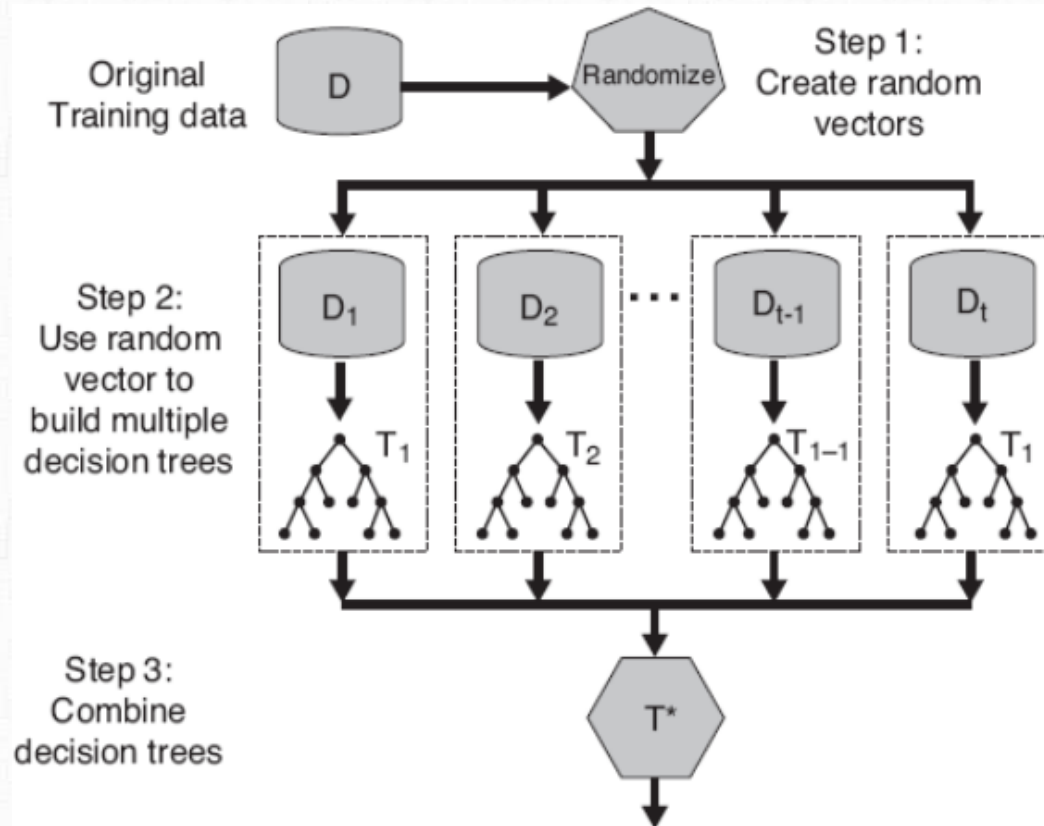
# Random Forests

---

- Ensemble method specifically designed for decision tree classifiers
- Introduce two sources of randomness:
  - “Bagging” and “Random input vectors”
  - **Bagging method**: each tree is grown using a bootstrap sample of training data
  - **Random vector method**: At each node, best split is chosen from a random sample of  $m$  attributes instead of all attributes



# Random Forests



```
library(randomForest)
```

```
set.seed(5123512)
```

```
fmodel <- randomForest(x=spamTrain[,spamVars],  
  y=spamTrain$spam,
```

```
  ntree=100,
```

```
# Note 4:  
#   Use 100 trees to be compatible with our bagging  
#   example. The default is 500 trees.
```

```
  nodesize=7,
```

```
# Note 5:  
#   Specify that each node of a tree must have a minimum  
#   of 7 elements, to be compatible with the default minimum node size that rpart()  
#   uses on this training set.
```

```
  importance=T)
```

```
# Note 6:  
#   Tell the algorithm to save information to be used for  
#   calculating variable importance (we'll see this later).
```

```
accuracyMeasures(predict(fmodel,  
  newdata=spamTrain[,spamVars],type='prob')[, 'spam'],  
  spamTrain$spam=="spam",name="random forest, train")
```

```
##                model  accuracy          f1  
## 1 random forest, train 0.9884142 0.9851943
```

```
accuracyMeasures(predict(fmodel,  
  newdata=spamTest[,spamVars],type='prob')[, 'spam'],  
  spamTest$spam=="spam",name="random forest, test")
```

```
##                model  accuracy          f1  
## 1 random forest, test 0.9497817 0.9340974
```

# Estimated Error Rate

---


1. At each bootstrap iteration, predict the data not in the bootstrap sample (what Breiman calls “out-of-bag”, or OOB, data) using the tree grown with the bootstrap sample.
2. Aggregate the OOB predictions. (On the average, each data point would be out-of-bag around 36% of the times, so aggregate these predictions.) Calculate the error rate, and call it the OOB estimate of error rate.

# Variable importance

---


- Make permutation on certain variable  $p$  of OOB data while other variables unchanged
- Compare prediction accuracy on permuted data and unchanged data

# Adaboost and XGBoost




## Random Forest

uses a set of fully grown trees




## Adaboost

uses a set of stumps  
– simplified decision tree (weak learner)




## Random Forest

all trees are equally important  
to make a final decision




## Adaboost

some stumps are more important  
than others  
when making a final decision




## Random Forest

makes all tree independently



## Adaboost


makes trees sequentially





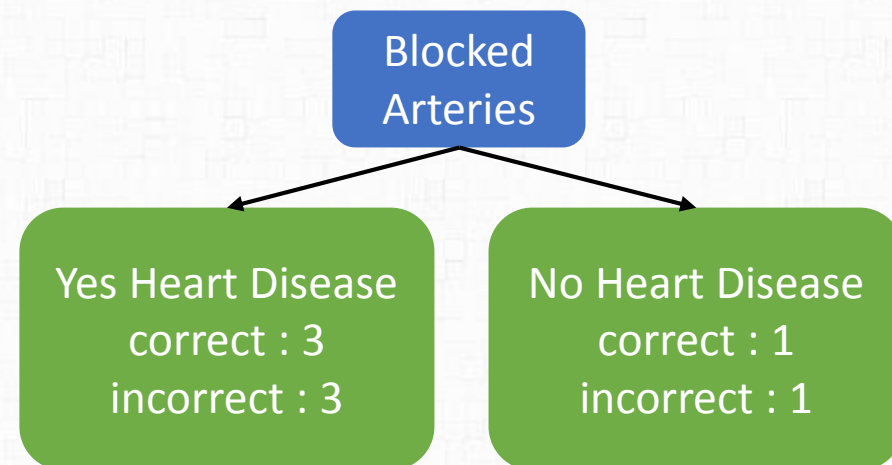

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8



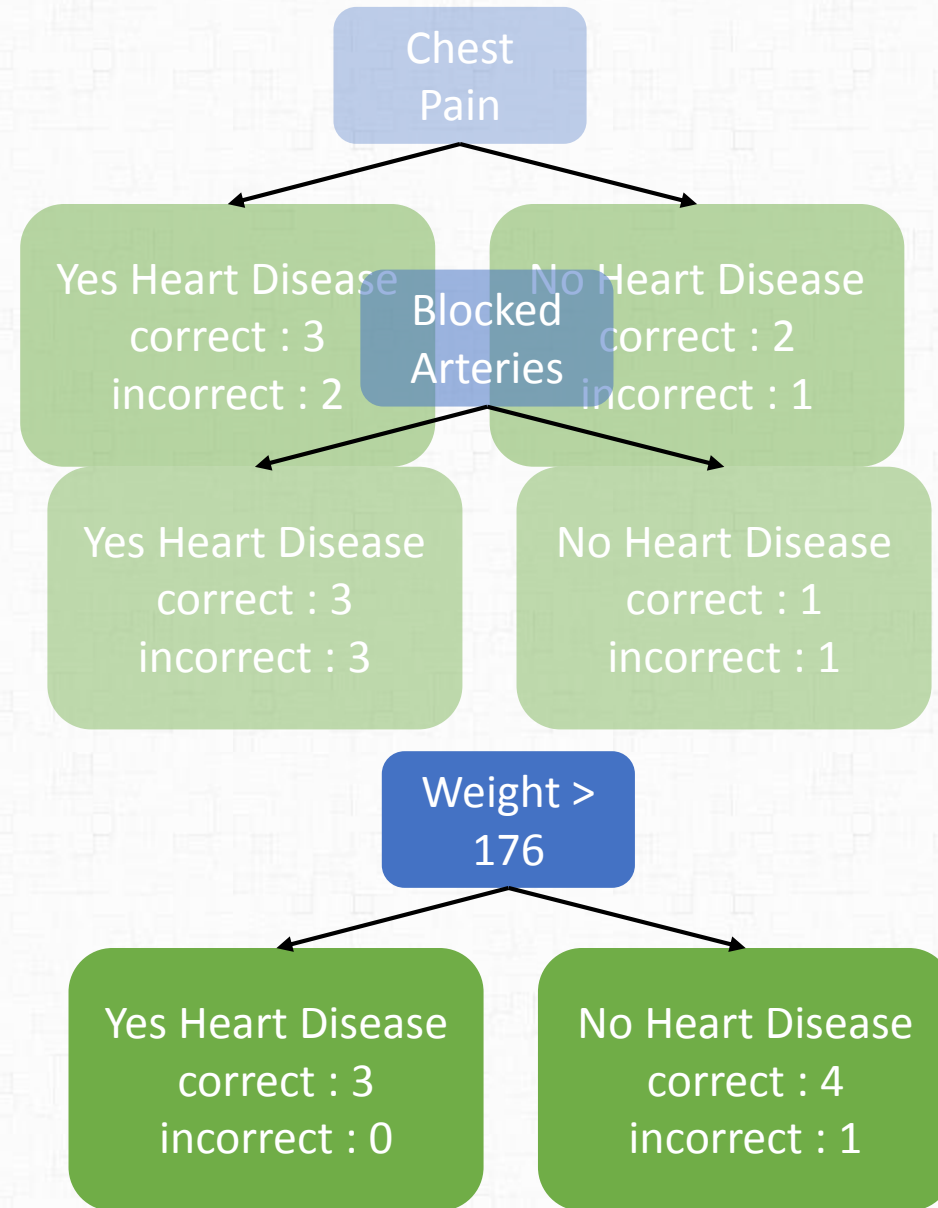
$$\begin{aligned}
 \text{G.I.} &= \frac{5}{8} \left[ 1 - \left\{ \left( \frac{3}{5} \right)^2 + \left( \frac{2}{5} \right)^2 \right\} \right] + \frac{3}{8} \left[ 1 - \left\{ \left( \frac{2}{3} \right)^2 + \left( \frac{1}{3} \right)^2 \right\} \right] \\
 &= 0.47
 \end{aligned}$$

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8



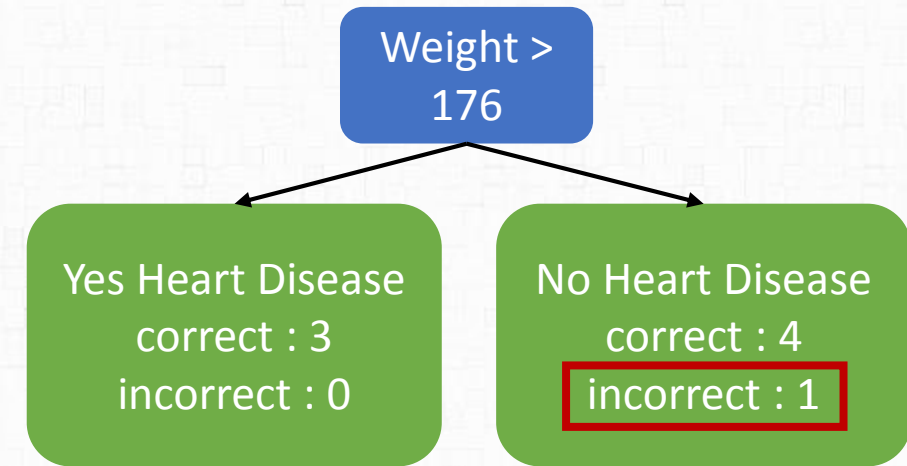
$$\begin{aligned}
 \text{G.I.} &= \frac{6}{8} \left[ 1 - \left\{ \left( \frac{3}{6} \right)^2 + \left( \frac{3}{6} \right)^2 \right\} \right] + \frac{2}{8} \left[ 1 - \left\{ \left( \frac{1}{2} \right)^2 + \left( \frac{1}{2} \right)^2 \right\} \right] \\
 &= 0.5
 \end{aligned}$$

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8



$$G.I. = \frac{3}{8} \left[ 1 - \left\{ \left( \frac{3}{3} \right)^2 + \left( \frac{0}{3} \right)^2 \right\} \right] + \frac{5}{8} \left[ 1 - \left\{ \left( \frac{4}{5} \right)^2 + \left( \frac{1}{5} \right)^2 \right\} \right] = 0.2$$

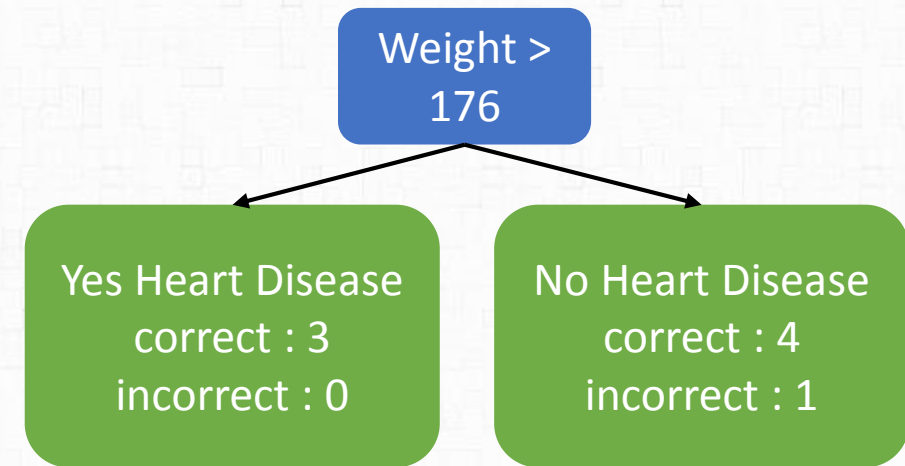

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8



Total Error: the sum of weights for *incorrectly* classified samples  
In the example, Total Error is 1/8

Stumps Weight:  
Amount of weight that the stump contributes to the final decision

$$\frac{1}{2} \log\left(\frac{1 - \text{total error}}{\text{total error}}\right)$$

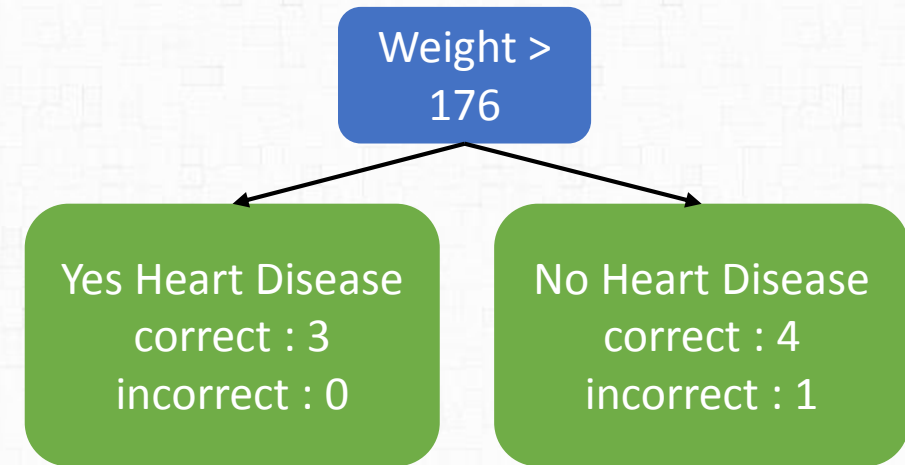



Total Error: the sum of weights for *incorrectly* classified samples  
In the example, Total Error is 1/8

Stumps Weight:

Amount of weight that the stump contributes to the final decision

$$\frac{1}{2} \log\left(\frac{1 - \text{total error}}{\text{total error}}\right)$$



Total Error: the sum of weights for *incorrectly* classified samples  
In the example, Total Error is  $\frac{1}{8}$

Stumps Weight:

Amount of weight that the stump contributes to the final decision

$$\frac{1}{2} \log \left( \frac{1 - \frac{1}{8}}{\frac{1}{8}} \right) = 0.97$$

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8


New Sample Weight:

Incorrectly classified samples :

$$\begin{aligned} & \text{sample weight} \times e^{\text{stump weight}} \\ &= \frac{1}{8} \times e^{0.97} = \frac{1}{8} \times 2.64 = 0.33 \end{aligned}$$

Correctly classified samples :

$$\text{sample weight} \times e^{\text{stump weight}}$$





Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8

New Sample Weight:

Incorrectly classified samples :


$$\text{sample weight} \times e^{\text{stump weight}}$$

$$= \frac{1}{8} \times e^{0.97} = \frac{1}{8} \times 2.64 = 0.33$$

Correctly classified samples :

$$\text{sample weight} \times e^{\text{stump weight}}$$

$$= \frac{1}{8} \times e^{-0.97} = \frac{1}{8} \times 0.38 = 0.05$$



Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight	Norm. Sample Weight
Yes	Yes	205	Yes	0.05	0.07
No	Yes	180	Yes	0.05	0.07
Yes	No	210	Yes	0.05	0.07
Yes	Yes	167	Yes	0.33	0.49
No	Yes	156	No	0.05	0.07
No	Yes	125	No	0.05	0.07
Yes	No	168	No	0.05	0.07
Yes	Yes	172	No	0.05	0.07

New Sample Weight:

Incorrectly classified samples :

$$\text{sample weight} \times e^{\text{stump weight}}$$

$$= \frac{1}{8} \times e^{0.97} = \frac{1}{8} \times 2.64 = 0.33$$

Correctly classified samples :

$$\text{sample weight} \times e^{\text{stump weight}}$$


$$= \frac{1}{8} \times e^{-0.97} = \frac{1}{8} \times 0.38 = 0.05$$

sum of new sample weights = 0.68

normalize the new weight dividing by 0.68

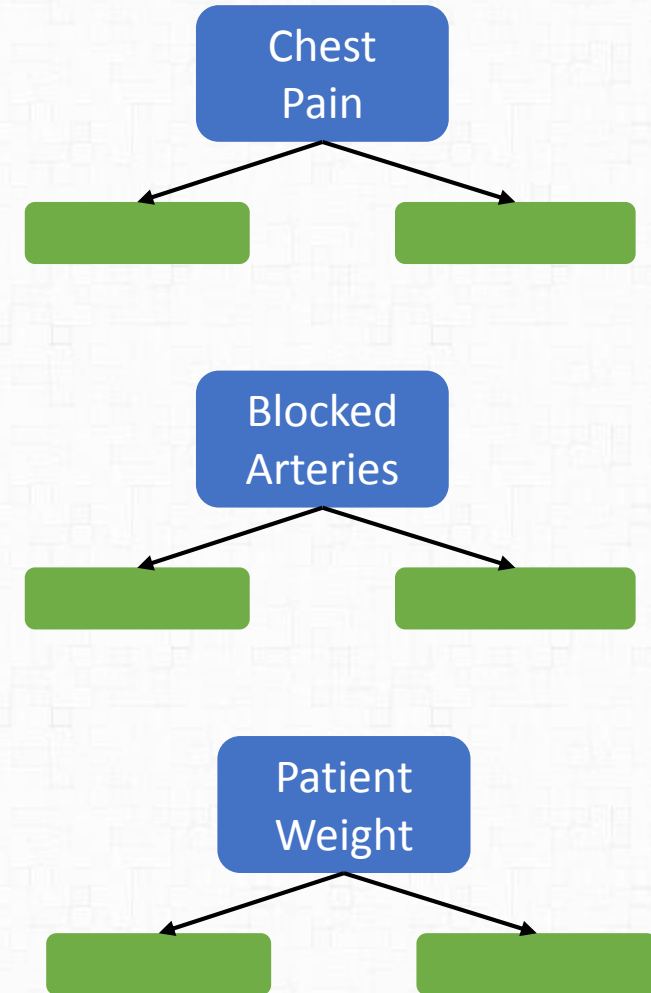
Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	0.07
No	Yes	180	Yes	0.07
Yes	No	210	Yes	0.07
<b>Yes</b>	<b>Yes</b>	<b>167</b>	<b>Yes</b>	0.49
No	Yes	156	No	0.07
No	Yes	125	No	0.07
Yes	No	168	No	0.07
Yes	Yes	172	No	0.07

make a new sample  
with the same sample size  
with respect to sample weight




Chest Pain	Blocked Arteries	Patient Weight	Heart Disease
No	Yes	156	No
<b>Yes</b>	<b>Yes</b>	<b>167</b>	<b>Yes</b>
No	Yes	125	No
<b>Yes</b>	<b>Yes</b>	<b>167</b>	<b>Yes</b>
<b>Yes</b>	<b>Yes</b>	<b>167</b>	<b>Yes</b>
Yes	Yes	172	No
<b>Yes</b>	<b>Yes</b>	<b>167</b>	<b>Yes</b>
Yes	Yes	167	Yes


Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
No	Yes	156	No	1/8
Yes	Yes	167	Yes	1/8
No	Yes	125	No	1/8
Yes	Yes	167	Yes	1/8
Yes	Yes	167	Yes	1/8
Yes	Yes	172	No	1/8
Yes	Yes	167	Yes	1/8
Yes	Yes	167	Yes	1/8




Now we have a forest of stumps...



## Heart Disease !




## No Heart Disease !



**Final Call**  
**The patient has Heart Disease !**


Heart Disease !



No Heart Disease !




>

**Total = 1.63**



# Gini Index (G. I.)


$$G.I(A) = \sum_{i=1}^d \left( R_i \left( 1 - \sum_{k=1}^m p_{ik}^2 \right) \right)$$






# XGBoost for Regression


Initial Prediction: 0.5



# XGBoost for Regression



# XGBoost for Regression



Initial Prediction: 0.5

-10.5, 6.5, 7.5, -7.5

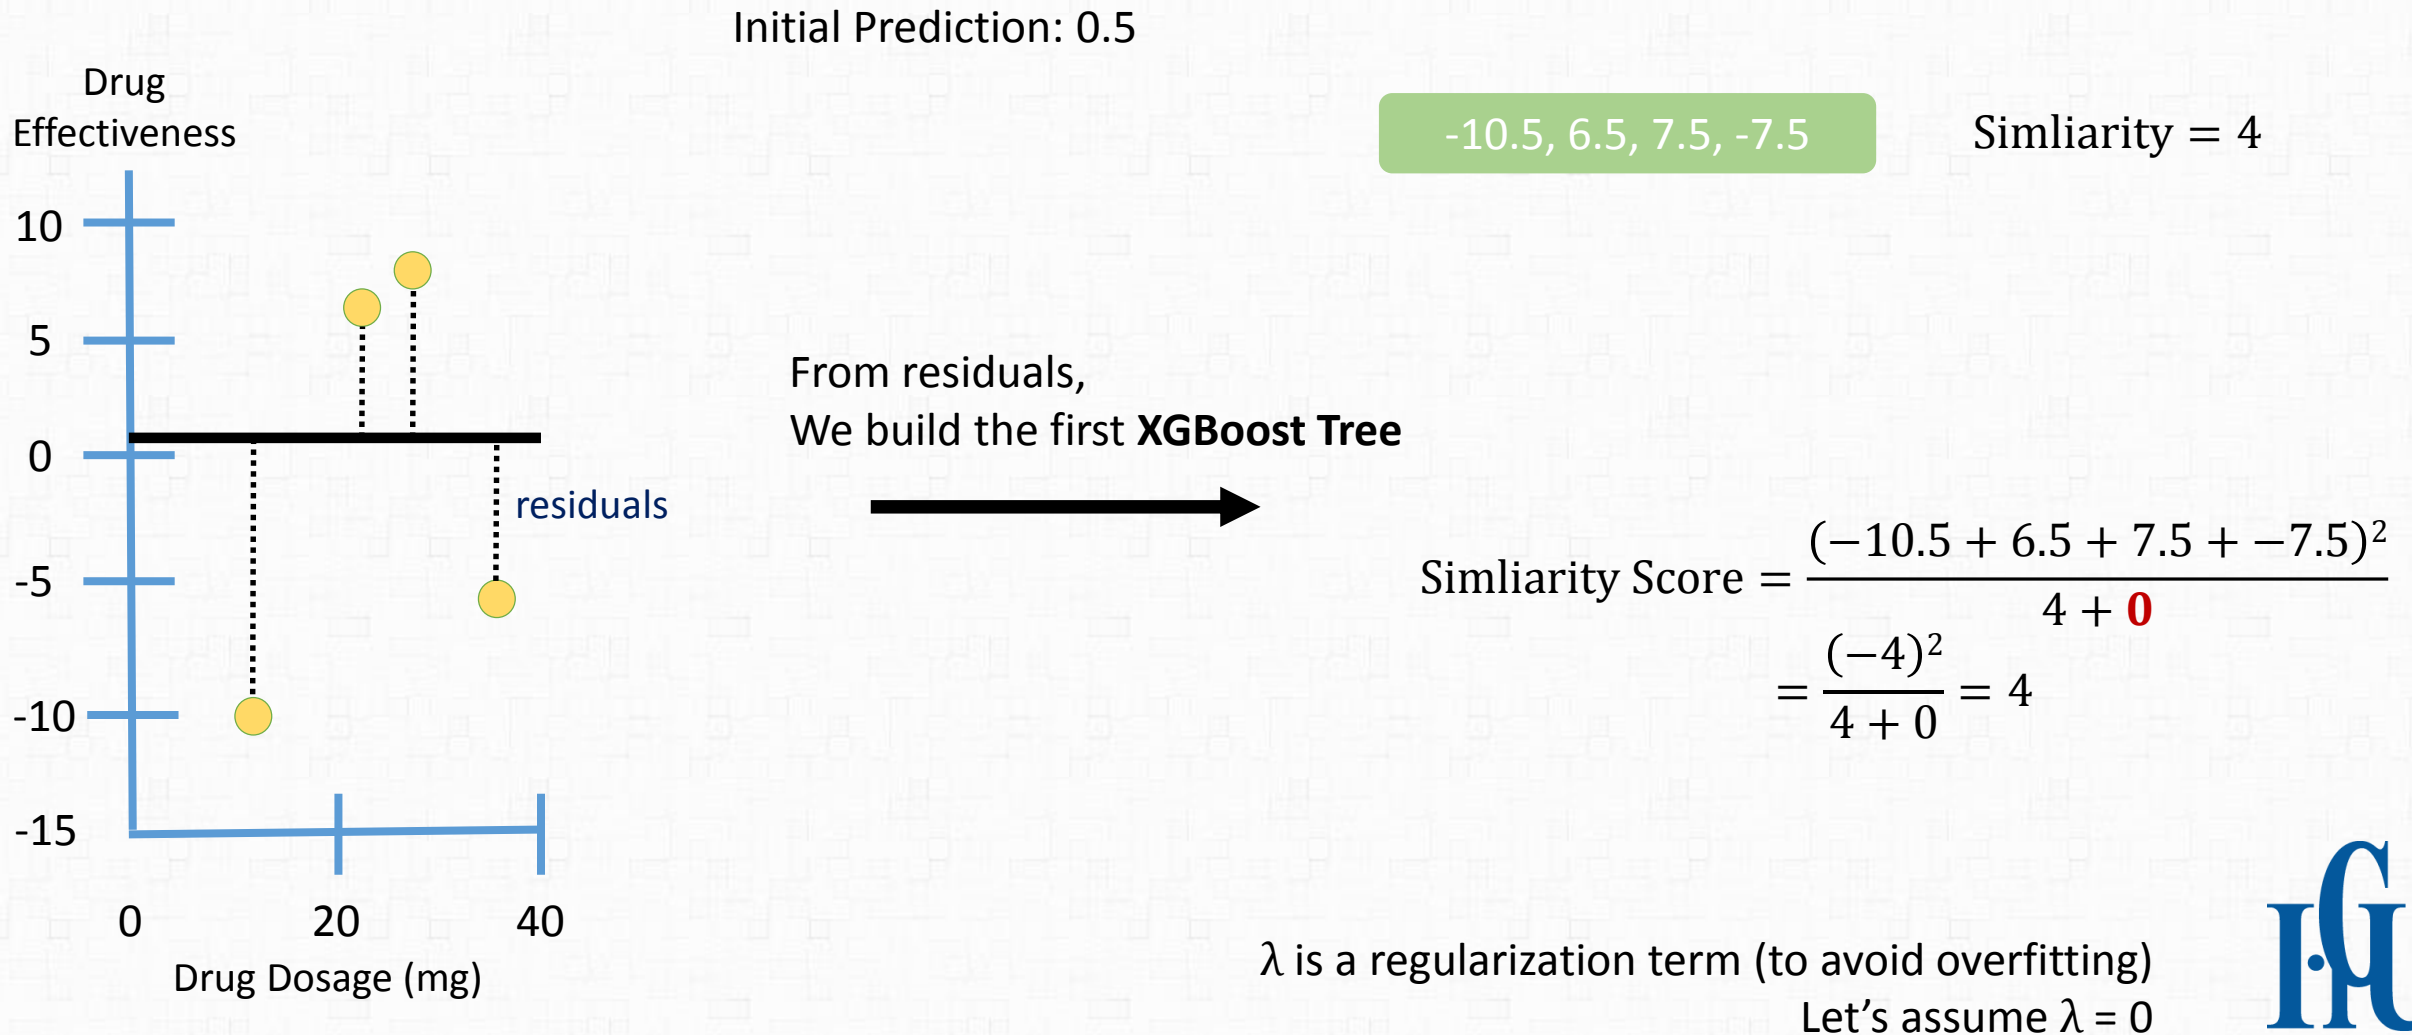
From residuals,  
We build the first **XGBoost Tree**




$$\text{Similarity Score} = \frac{\text{Sum of Residual, Squared}}{\text{Number of Residuals} + \lambda(\text{lambda})}$$

$\lambda$  is a regularization term (to avoid overfitting)  
Let's assume  $\lambda = 0$


# XGBoost for Regression



# XGBoost for Regression



Initial Prediction: 0.5

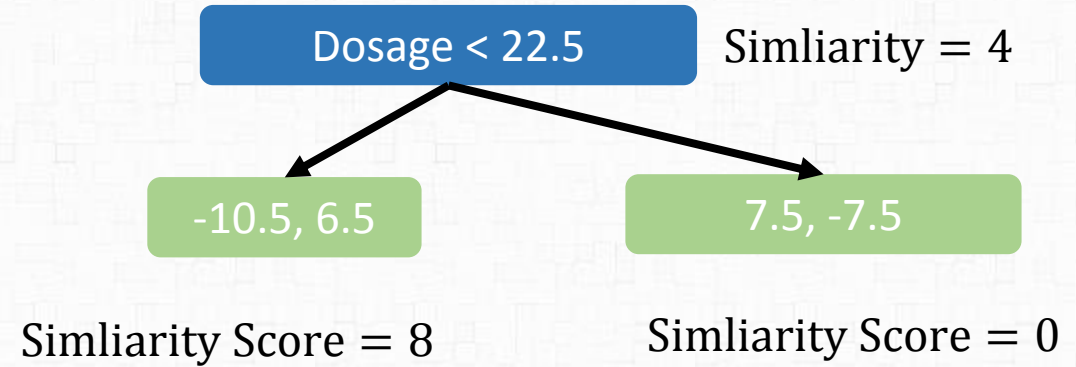



$$\text{Simliarity Score} = \frac{(-10.5)^2}{1 + \mathbf{0}} = 110.25$$

$$\text{Simliarity Score} = \frac{(6.5 + 7.5 + -7.5)^2}{3 + \mathbf{0}} = 14.08$$

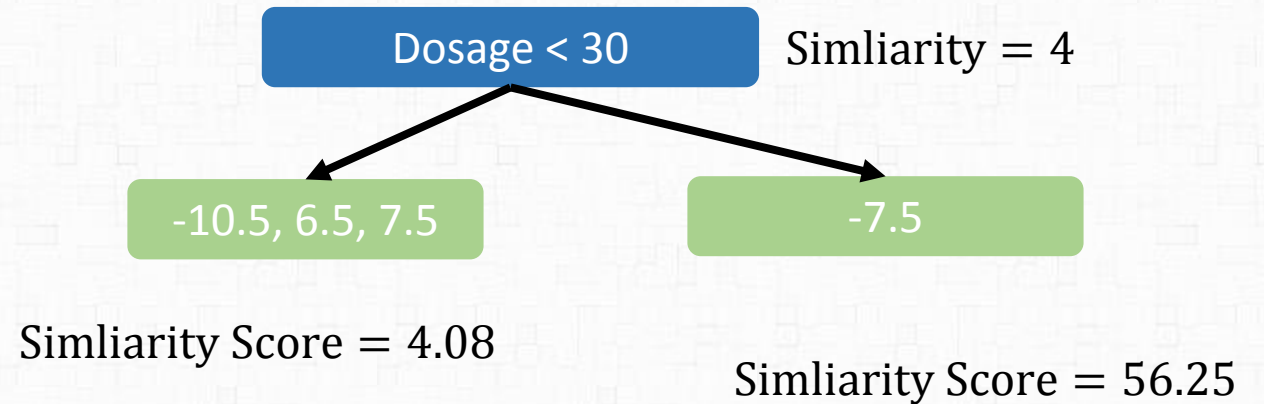

$$\text{Gain} = \text{Leftsimiliarity} + \text{Rightsimiliarity} - \text{Rootsimiliarity} = 120.33$$

# XGBoost for Regression



$$\text{Gain} = 8 + 0 - 4 = 4$$

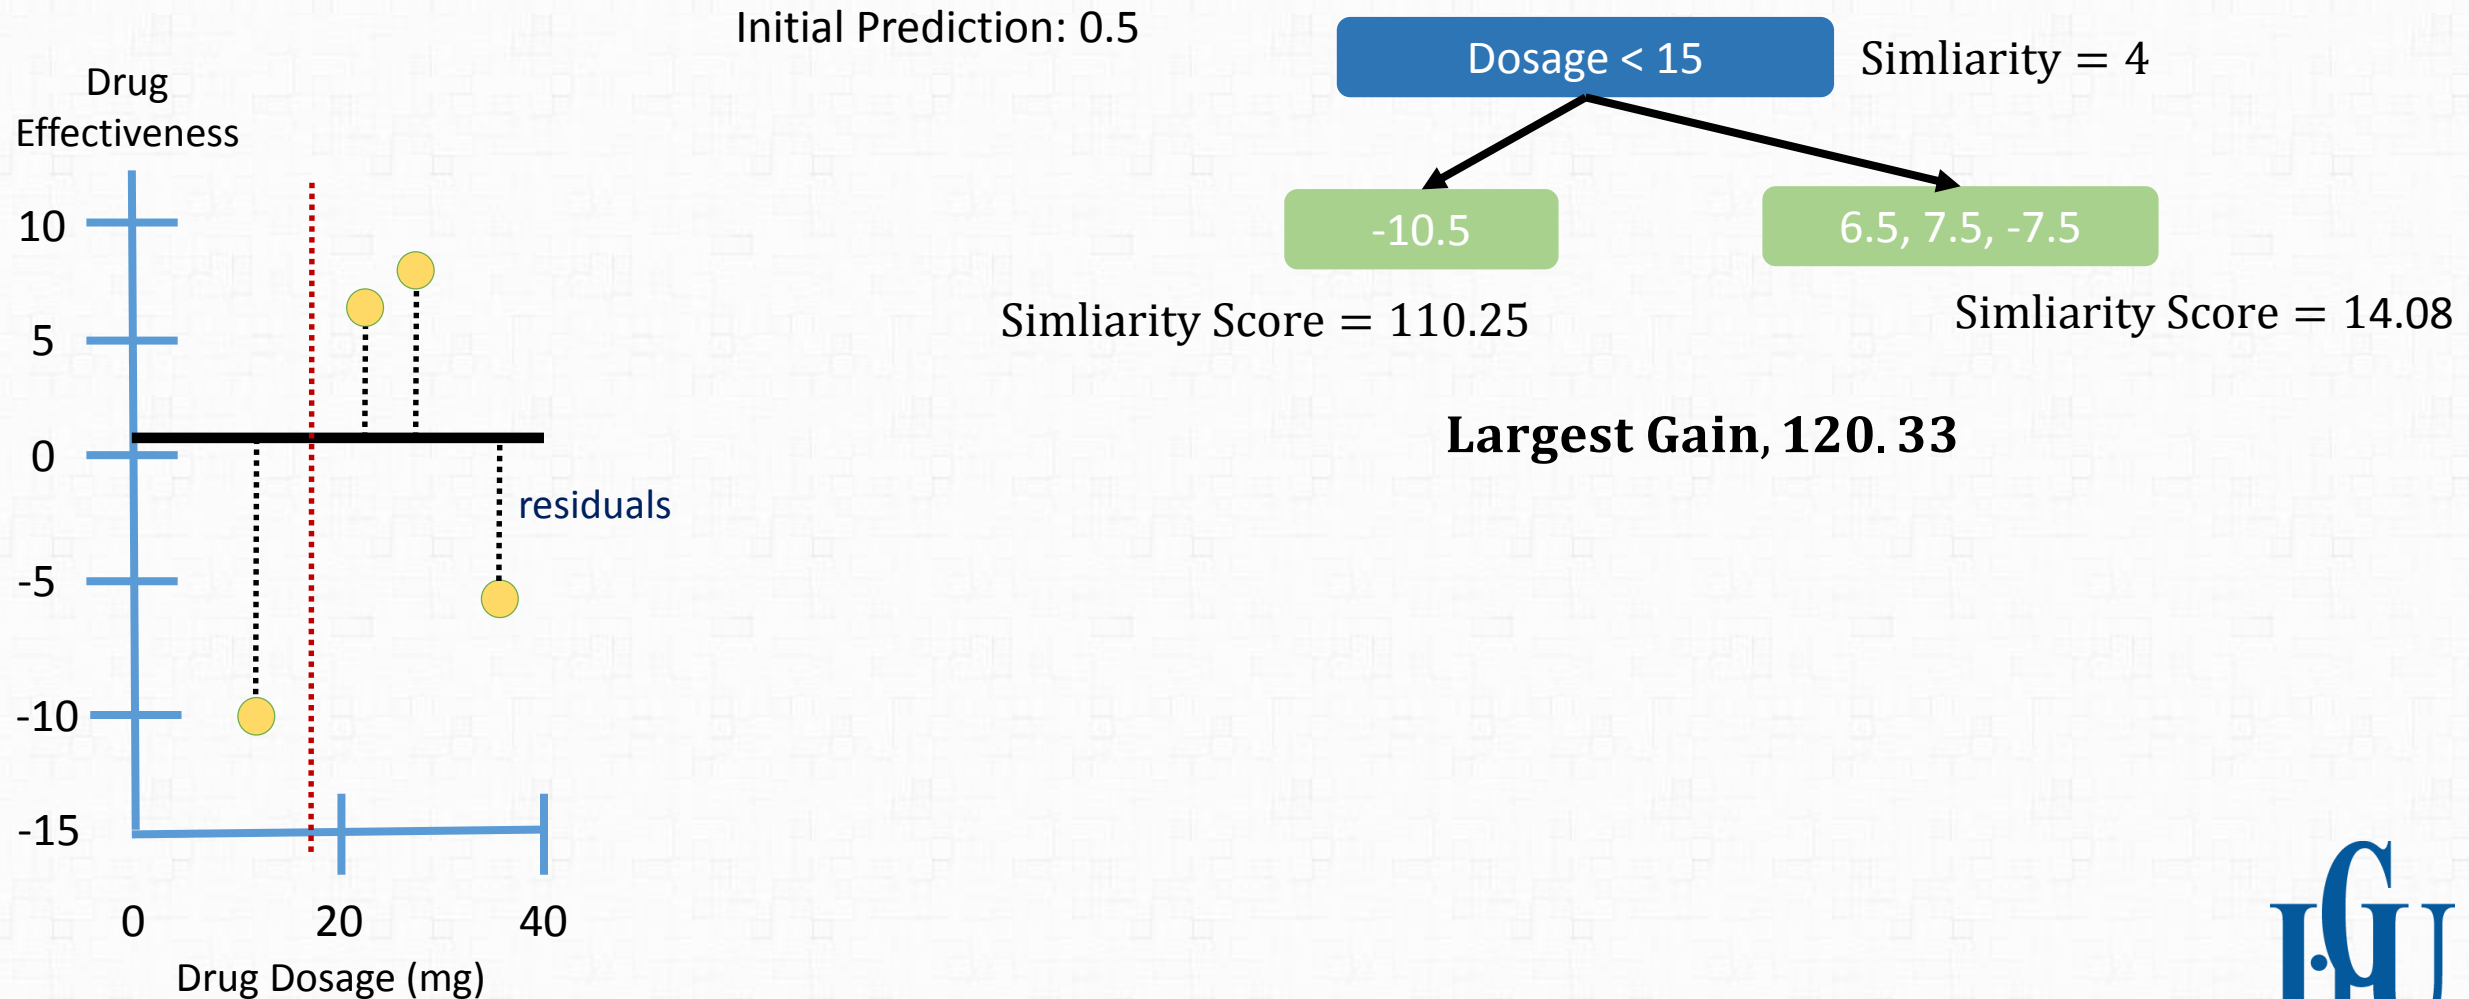
# XGBoost for Regression



$$\text{Gain} = \text{Leftsimilarity} + \text{Rightsimilarity} - \text{Rootsimilarity} = 56.33$$

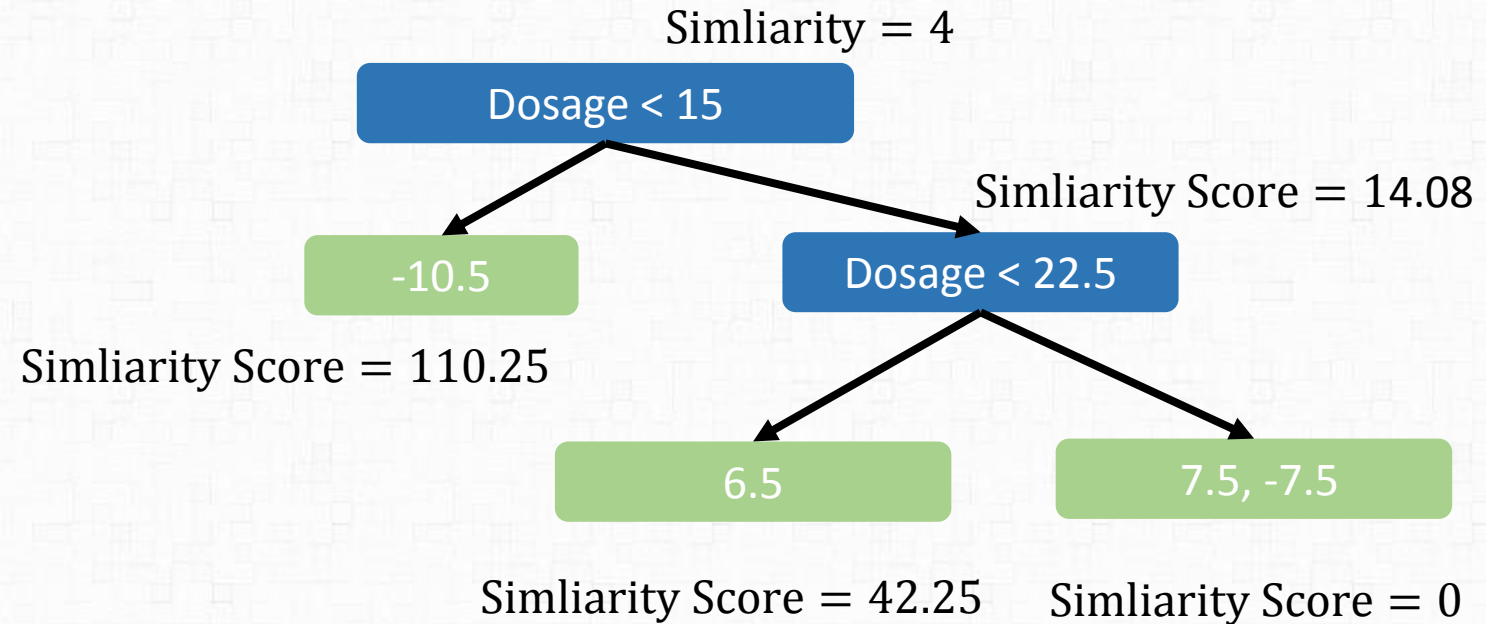



# XGBoost for Regression





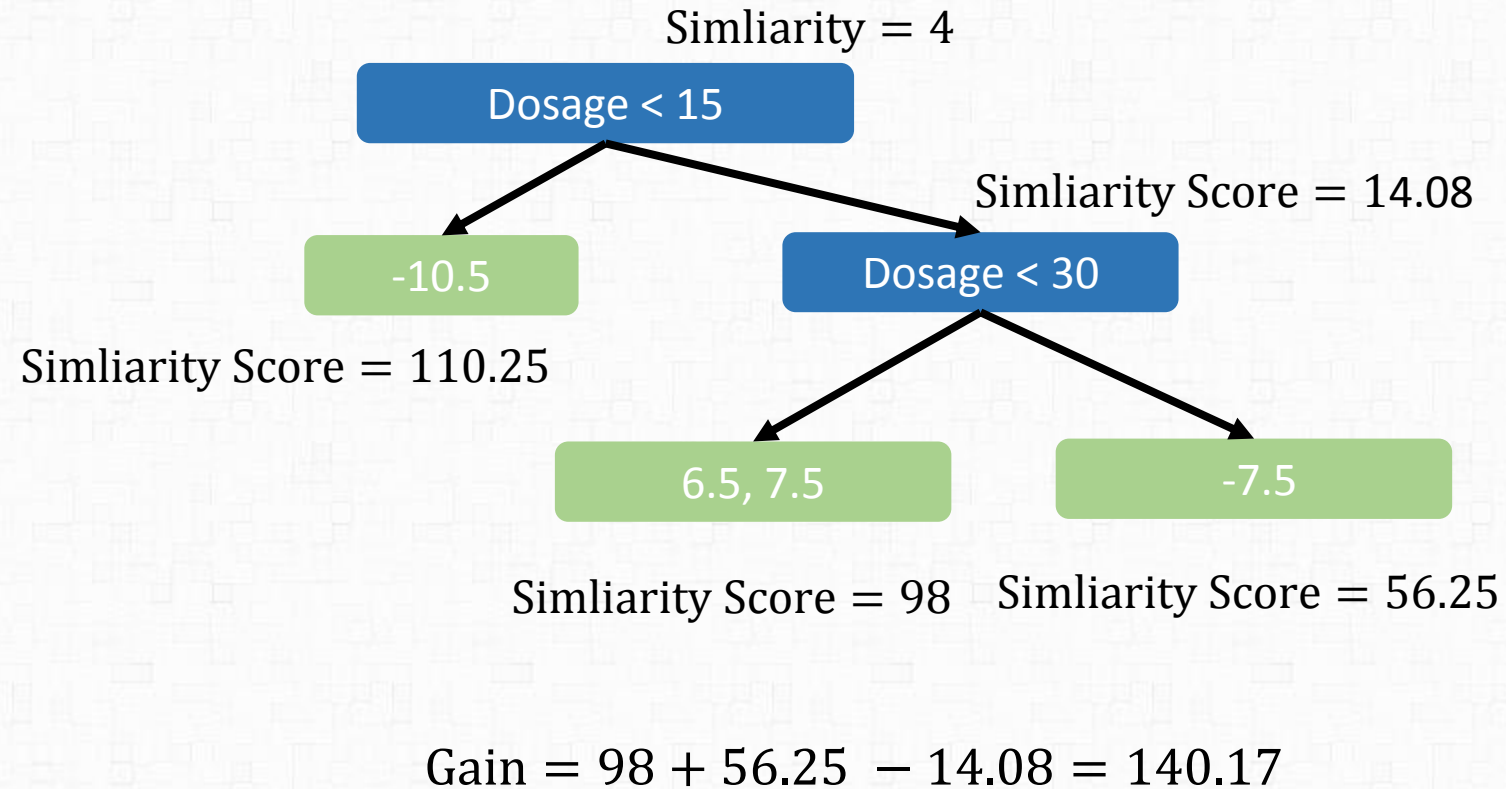

# XGBoost for Regression



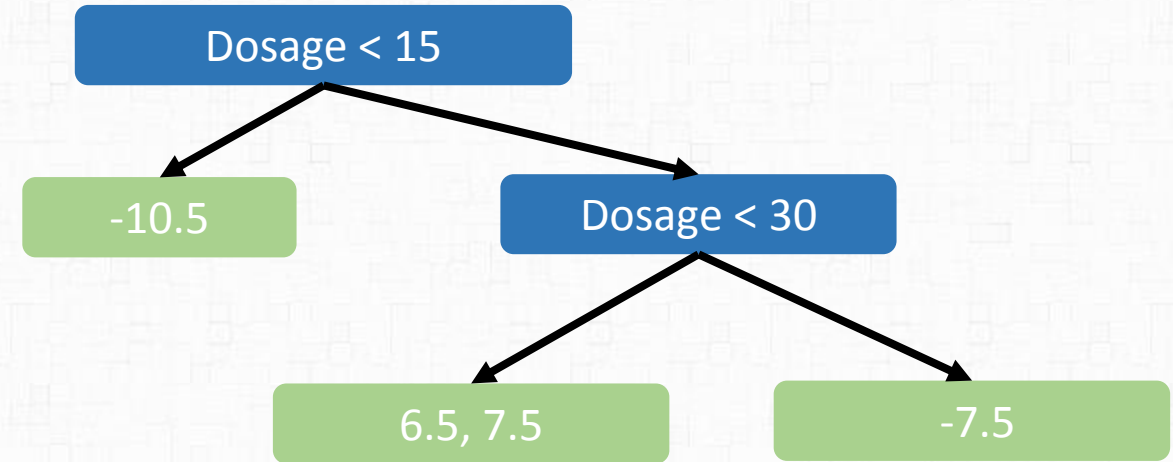

$$\text{Gain} = \text{Leftsimilarity} + \text{Rightsimilarity} - \text{Rootsimilarity} =$$

$$42.25 + 0 - 14.08 = 28.17$$

# XGBoost for Regression



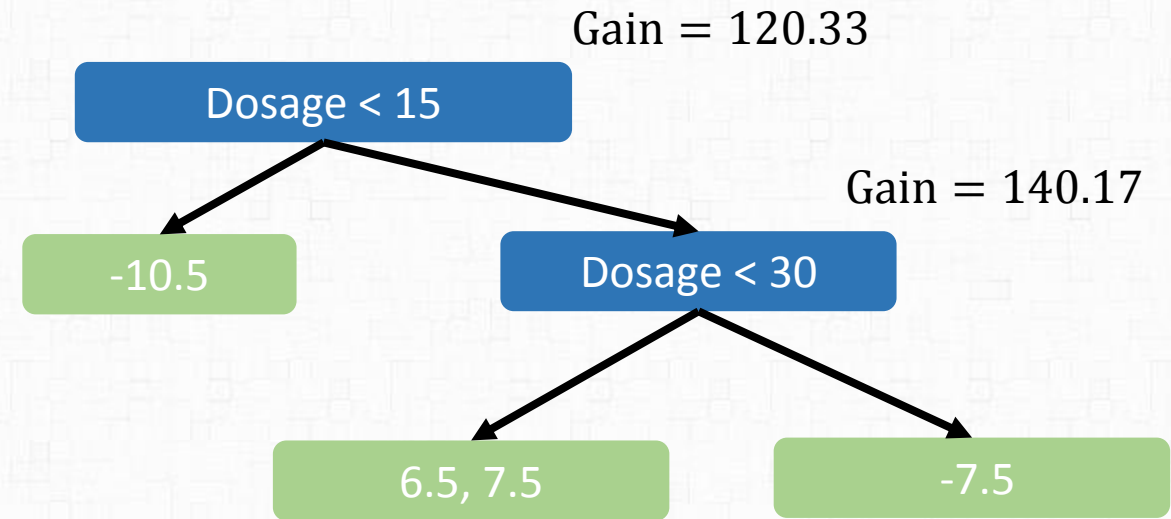

# XGBoost for Regression



Limiting max depth of XGBoost Tree as 2,  
We stop branching out the tree here.

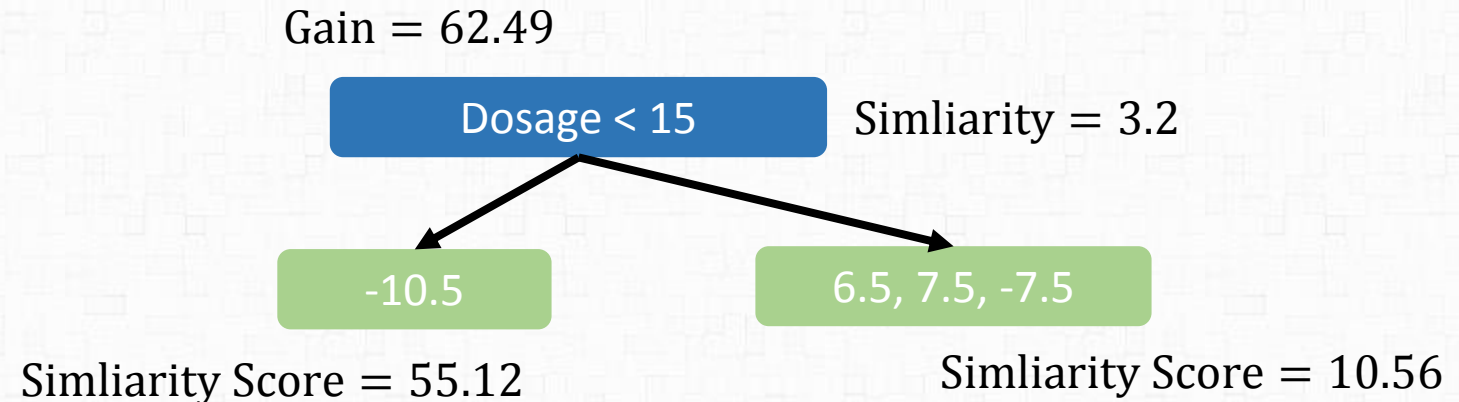

By default, max depth is 6 levels

# XGBoost for Regression



We prune the tree comparing Gain and  $\gamma$  (gamma) (starting from leaf nodes)  
 $\gamma$  is threshold for Gain  
 when the max gain is smaller than  $\gamma$ , we prune that branch

# Regularization term - $\lambda$



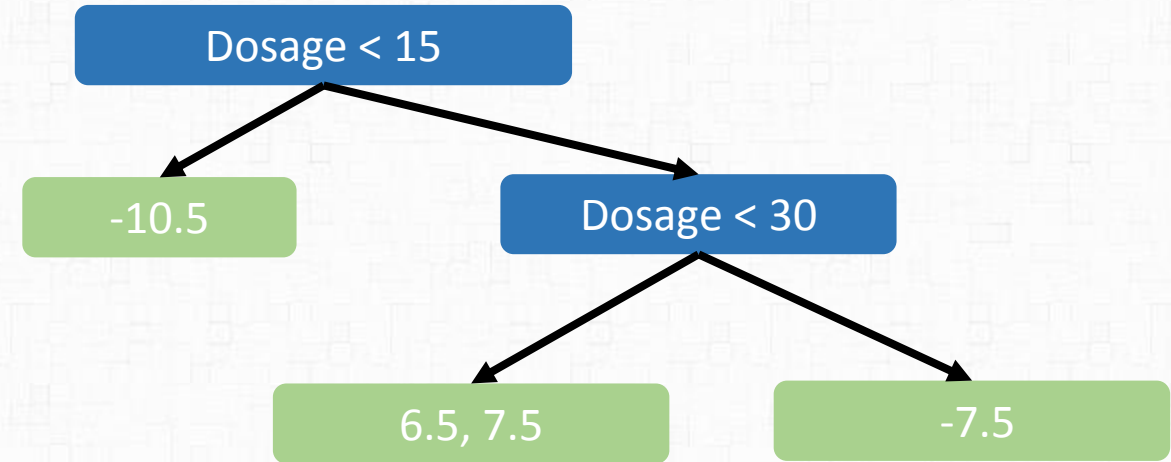

**Now let's think about  $\lambda > 0$**

**Let's assume  $\lambda = 1$**

For the smaller number of residuals,  
it decrease more of similarity score

Since it gives smaller Gain,  
Tree gets pruned more easily

# XGBoost for Regression

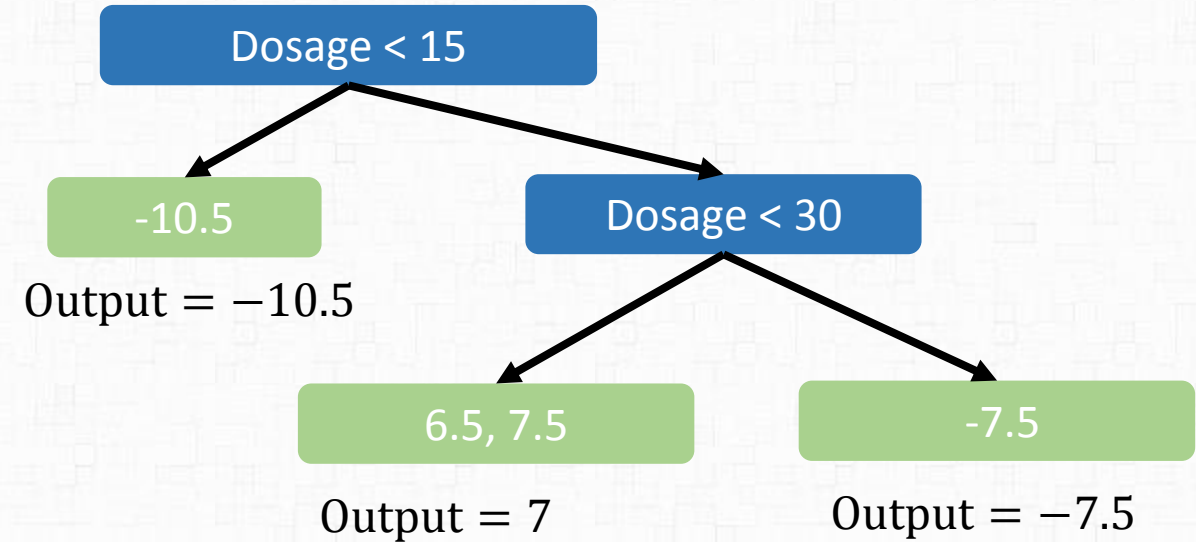



$$\text{Output Value} = \frac{\text{Sum of Residual}}{\text{Number of Residuals} + \lambda}$$

$$\text{Simliarity Score} = \frac{\text{Sum of Residual, Squared}}{\text{Number of Residuals} + \lambda}$$



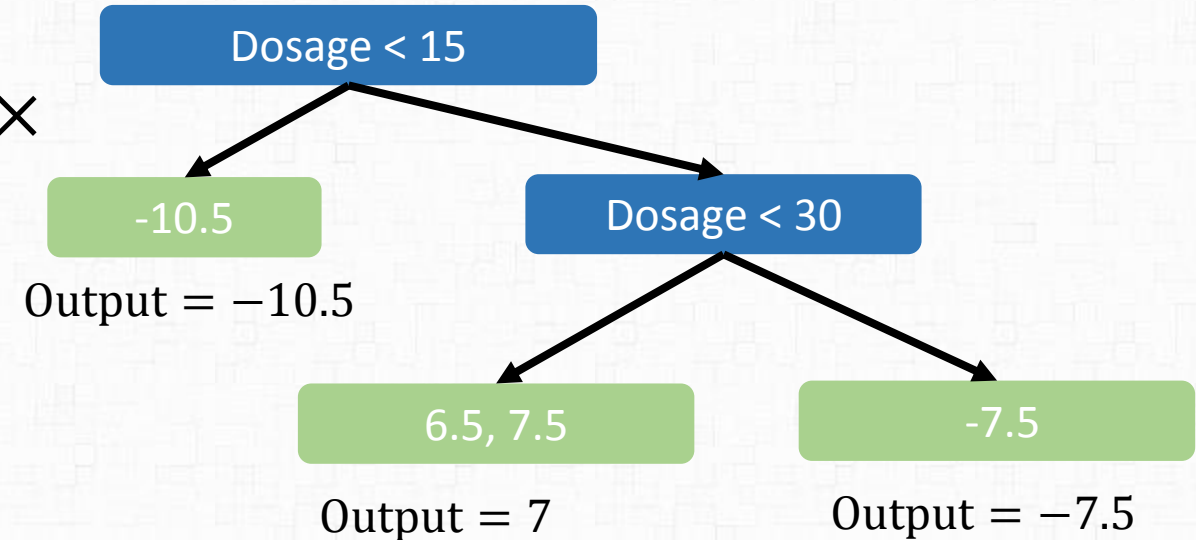

# XGBoost for Regression



$$\text{Output Value} = \frac{\text{Sum of Residual}}{\text{Number of Residuals} + \lambda}$$


$$\text{Simliarity Score} = \frac{\text{Sum of Residual, Squared}}{\text{Number of Residuals} + \lambda}$$

# XGBoost for Regression

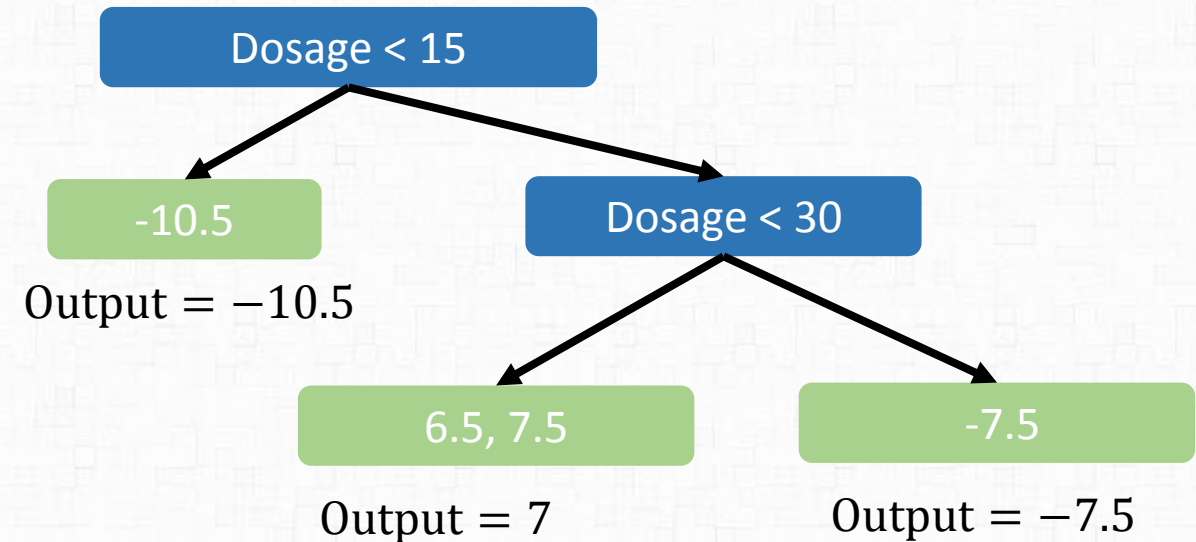




# XGBoost for Regression

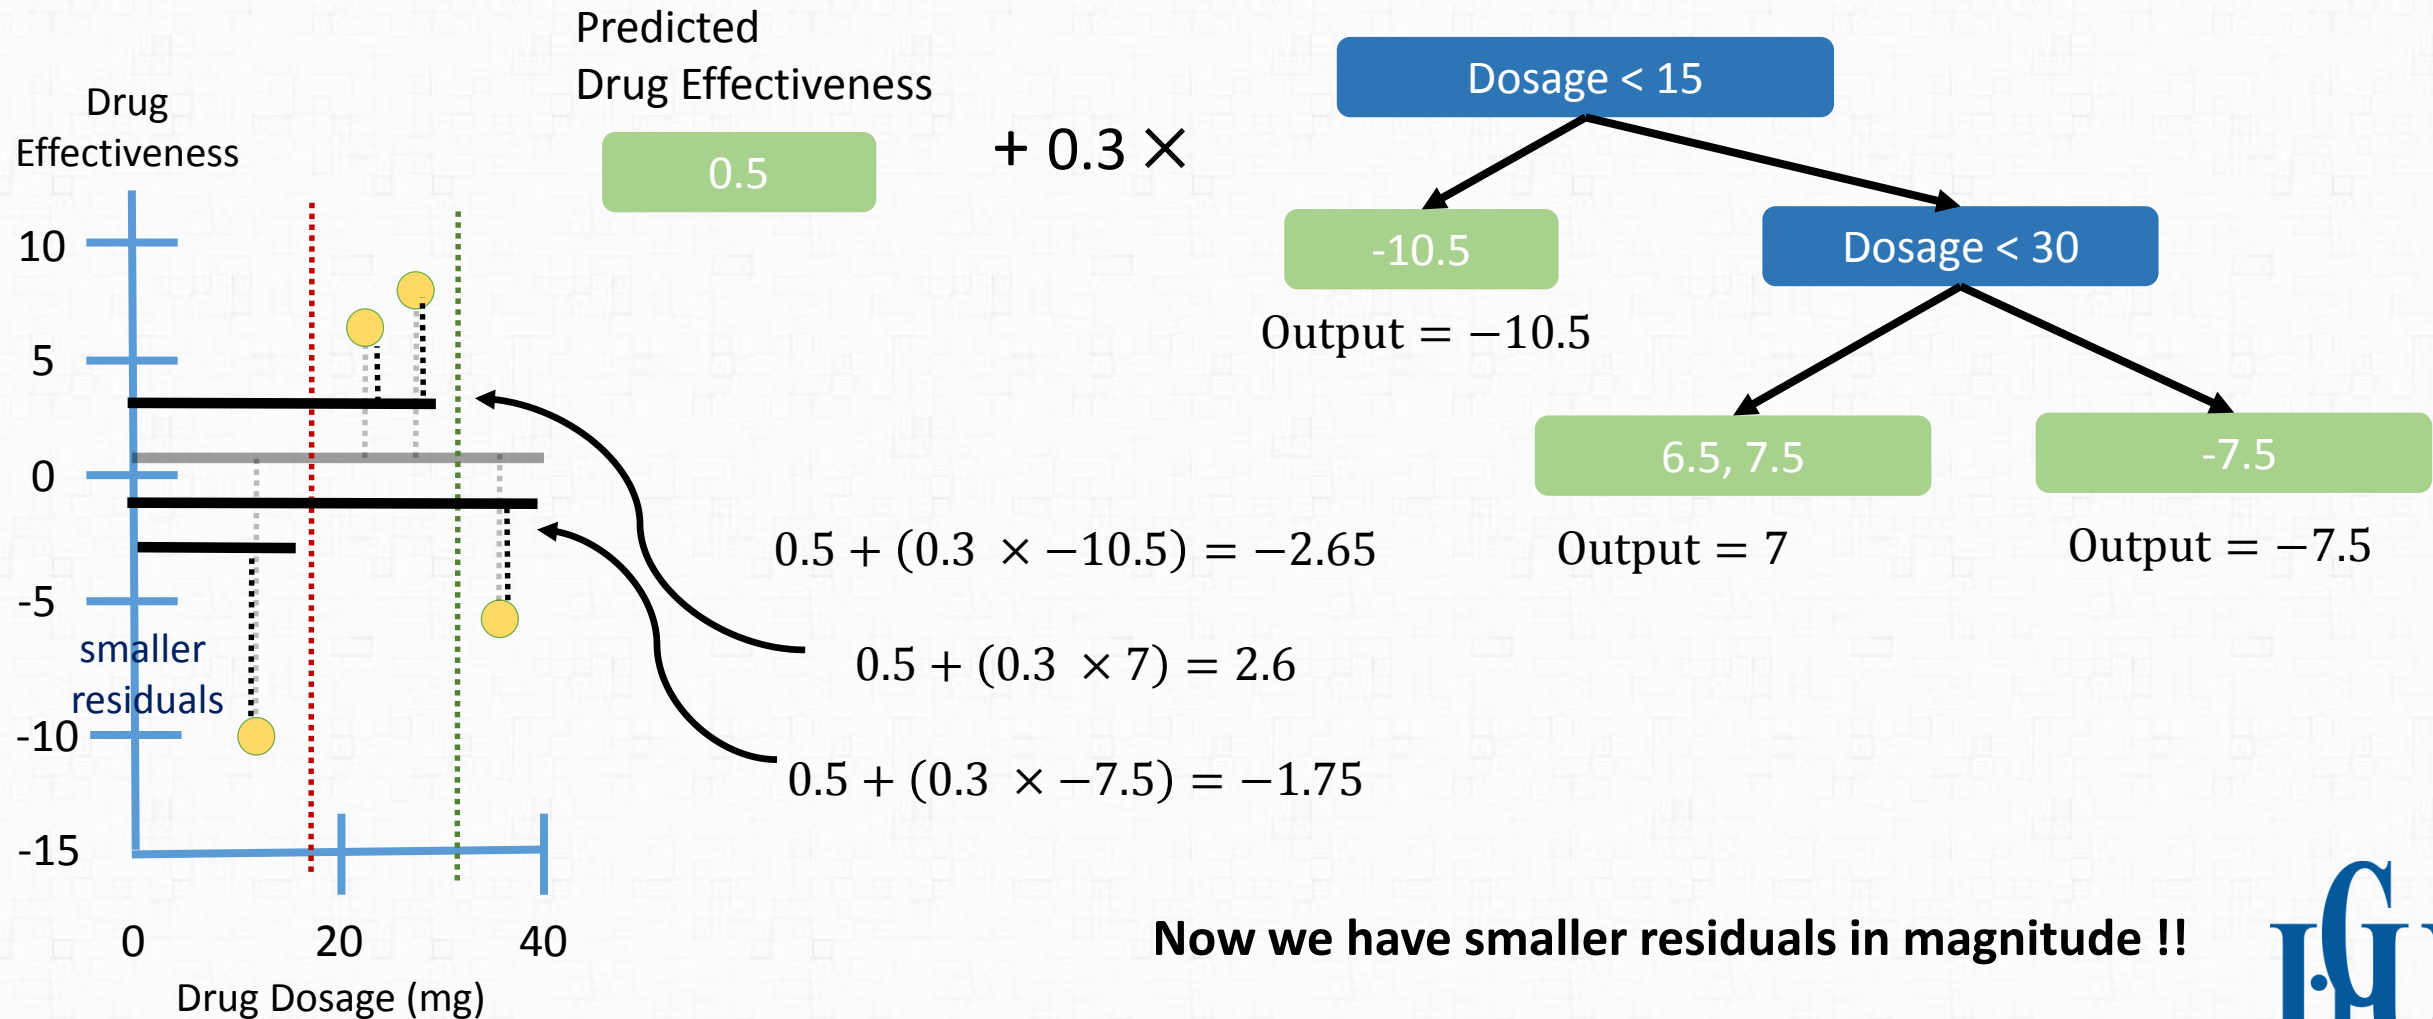


New predicted value is  
 $0.5 + (0.3 \times -10.5) = -2.65$

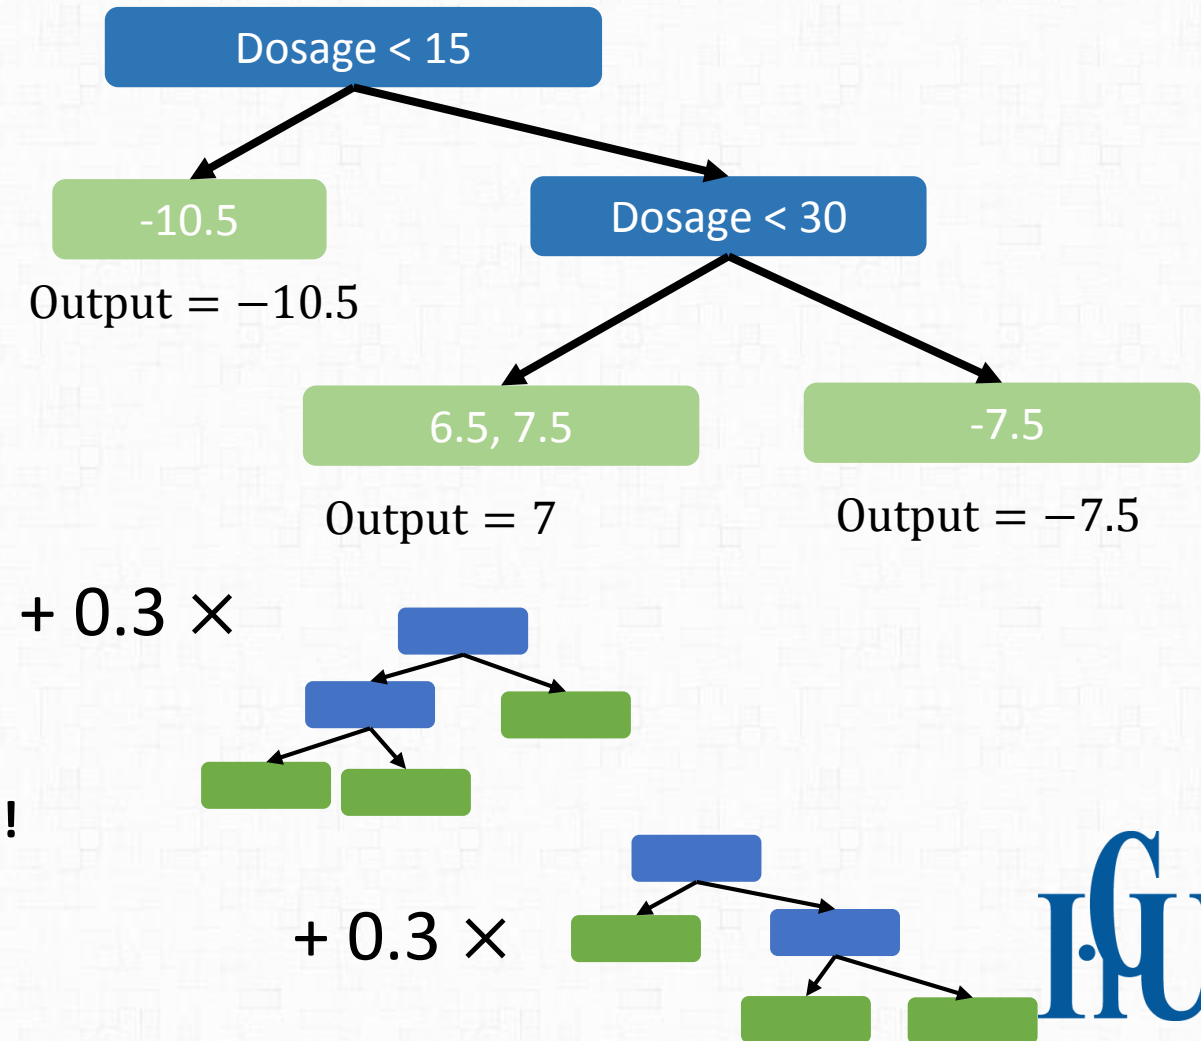



Learning rate is called  $\epsilon$ (eta),  
 by default 0.3 In XGBoost

# XGBoost for Regression



# XGBoost for Regression



# What to do next

---

- XGBoost for Classification, <https://youtu.be/8b1JEDvenQU>
- Mathematical details, <https://youtu.be/ZVFeW798-2I>
- Properties, <https://youtu.be/oRrKeUCEbq8>

# References

---

- STATQUEST, <https://statquest.org/>